

Explore Transformer-Based Foundation Models for Time Series Predictive Maintenance

Elene Esakia

BSc Computer Science

Constructor University

Supervisor: Hendro Wicaksono, Rahma Dawud

August 2025

Content

1.. Introduction

5. Conclusion

2. Relevant Work

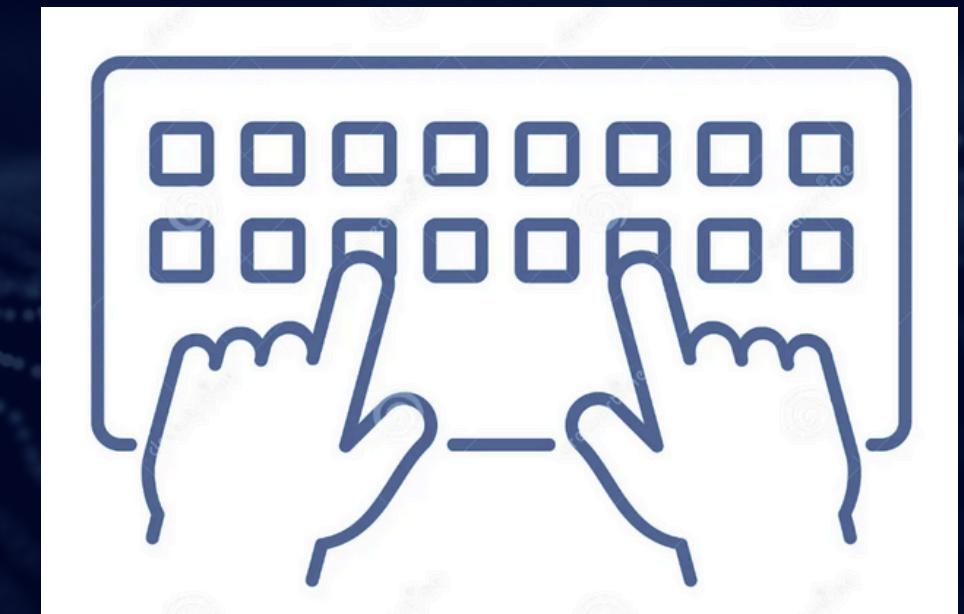
6. Future Work

3. Methodology

7. Q&A

4. Results & Analysis

8. References



Introduction

1. Project goals:

- Investigate transformer-based models for time series predictive maintenance
- Compare their performance against traditional approaches

2. Importance in Industry:

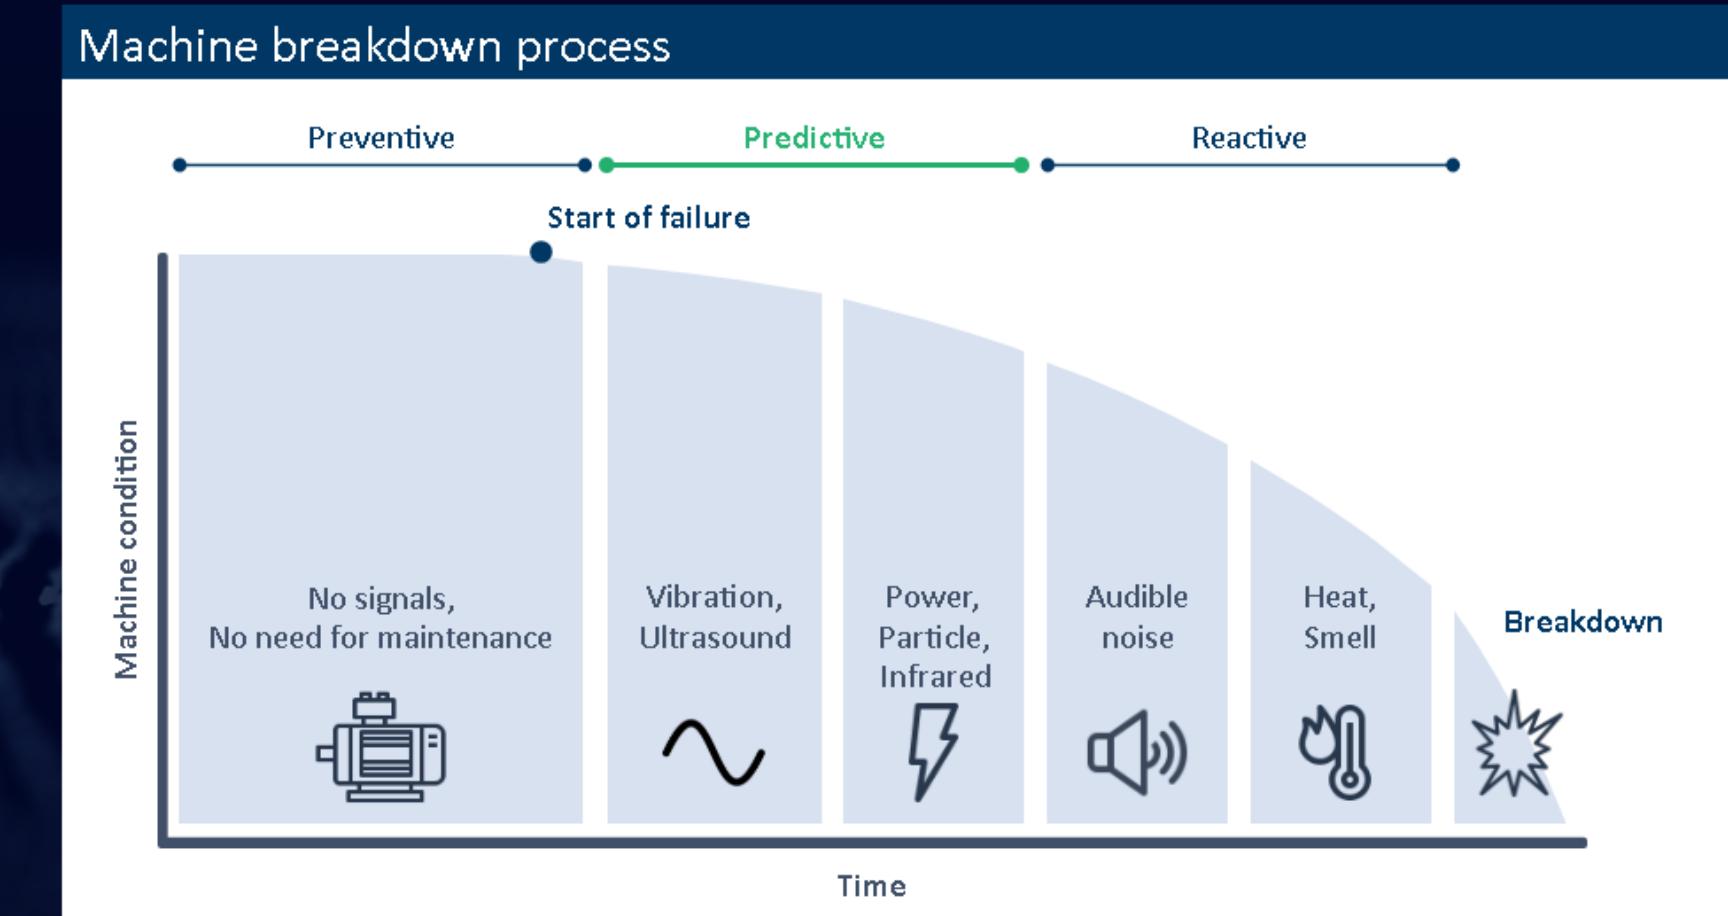
- Prevent unexpected machine breakdowns
- Save time and money on repairs
- Make machines more reliable and safe

3. Why Transformers:

- Learn patterns from time series data
- Handle multiple sensor readings at once

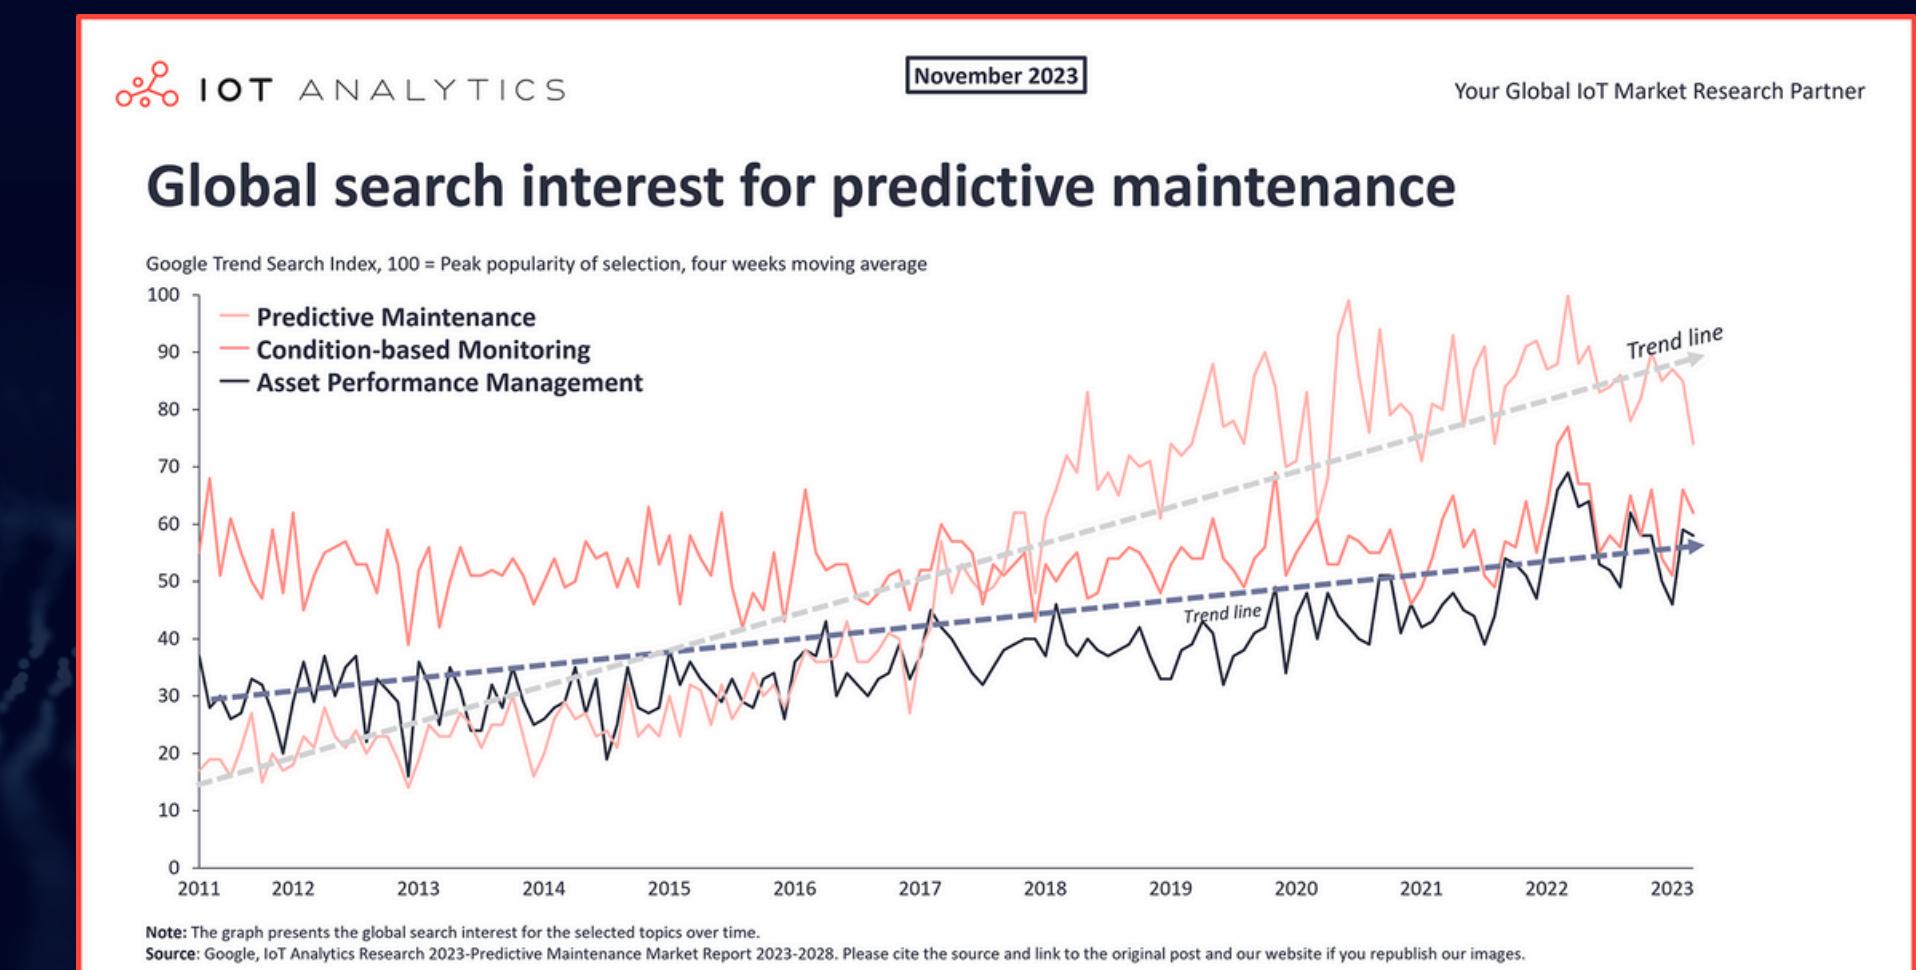
4. Models Explored:

- PatchTST
- MOMENT



Relevant Work

- Predictive maintenance market valued at \$5.5B in 2022 (+11% from 2021)
- Unplanned downtime costs ~\$125K per hour across industries
- One accurate prediction can save >\$100K
- Challenges: Many solutions have <50% accuracy, causing false alarms
- Research gaps:
 - Limited research on transformers for real-world PdM classification
 - Challenge: imbalanced datasets, rare failures, high computational cost
 - Unclear if these models can outperform traditional approaches in practical PdM scenarios



Methodology

Step 1: Data Exploration (EDA)

1. Dataset:

- Microsoft Azure Predictive Maintenance dataset (available on Kaggle)
- Captures 1 year of hourly operational data for 100 industrial machines
- 5 CSV files: PdM_telemetry.csv, PdM_failures.csv, PdM_errors.csv, PdM_maint.csv, PdM_machines.csv

2. Key Statistics:

- 876,100 records
- Total failures: 761 (98 machines affected)
- Sensors: 4 features (volt, rotate, pressure, vibration)
- Missing values: none; duplicates: none
- Prediction windows: 24h / 48h / 72h (Use Case)

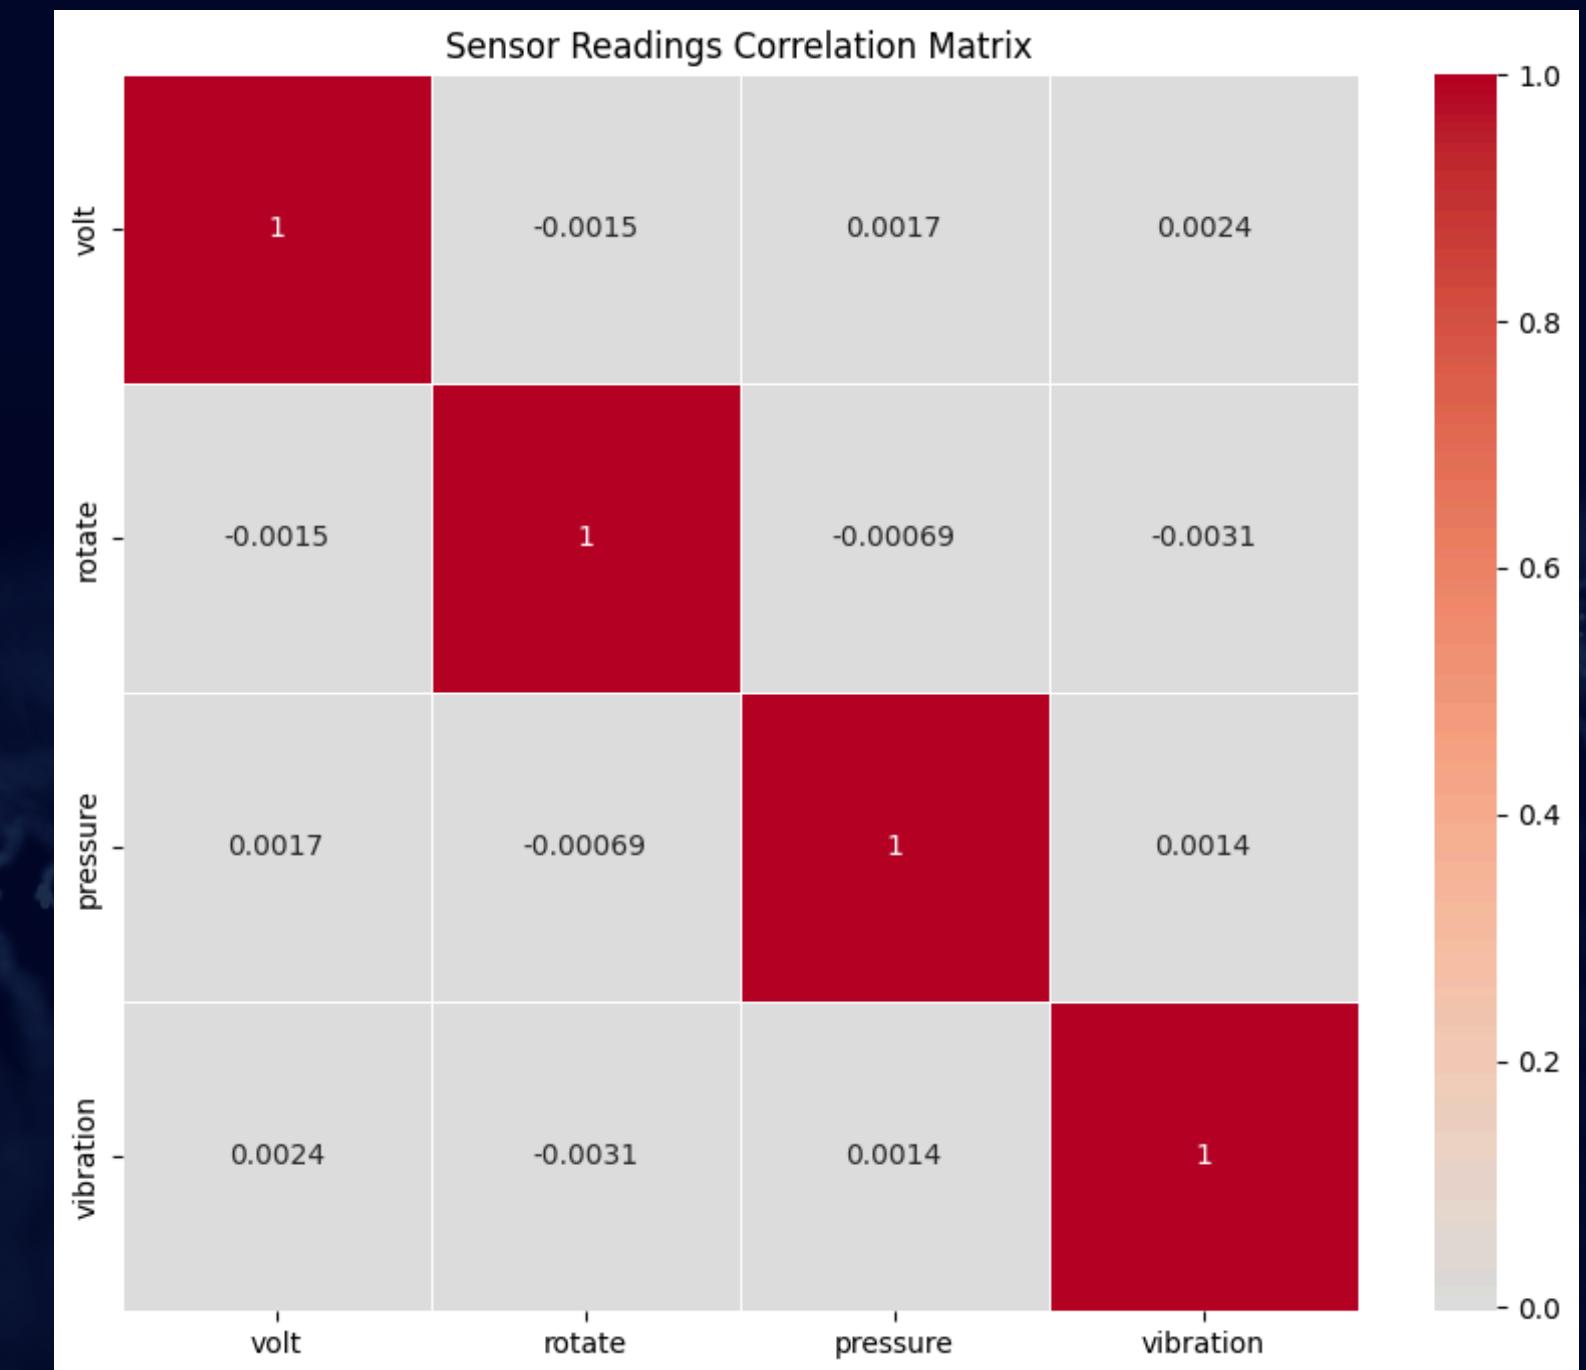


Fig. 1: Sensor Correlation Matrix (telemetry)

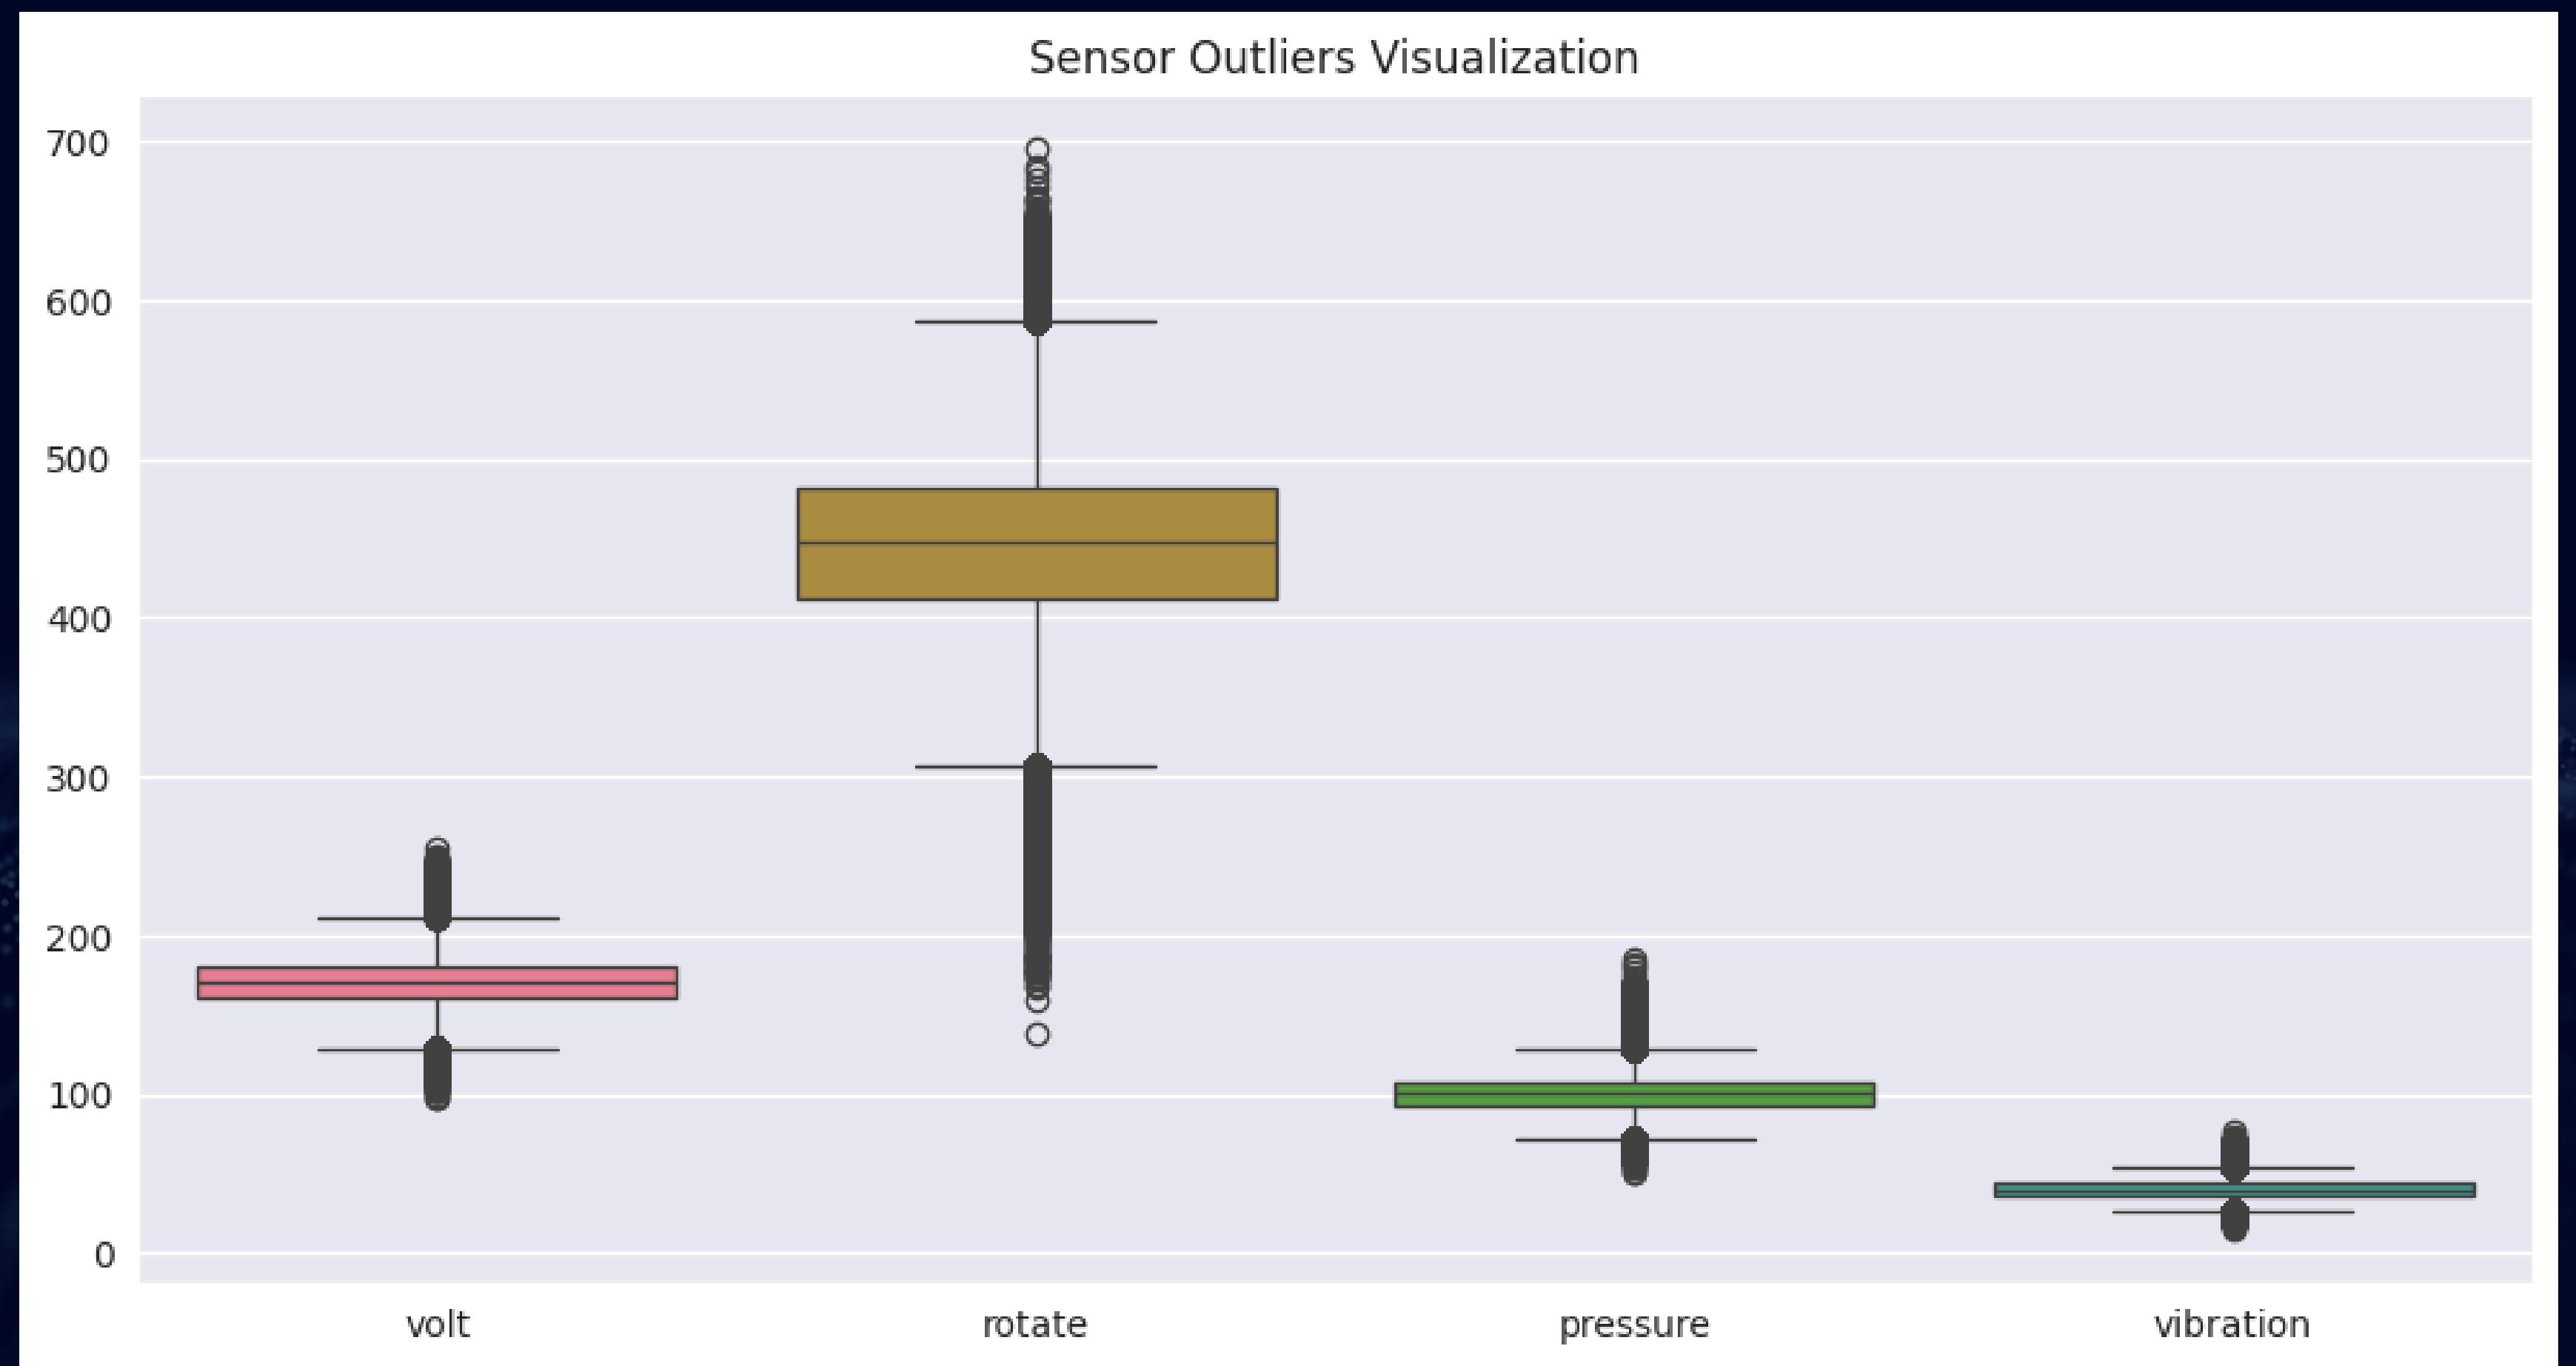
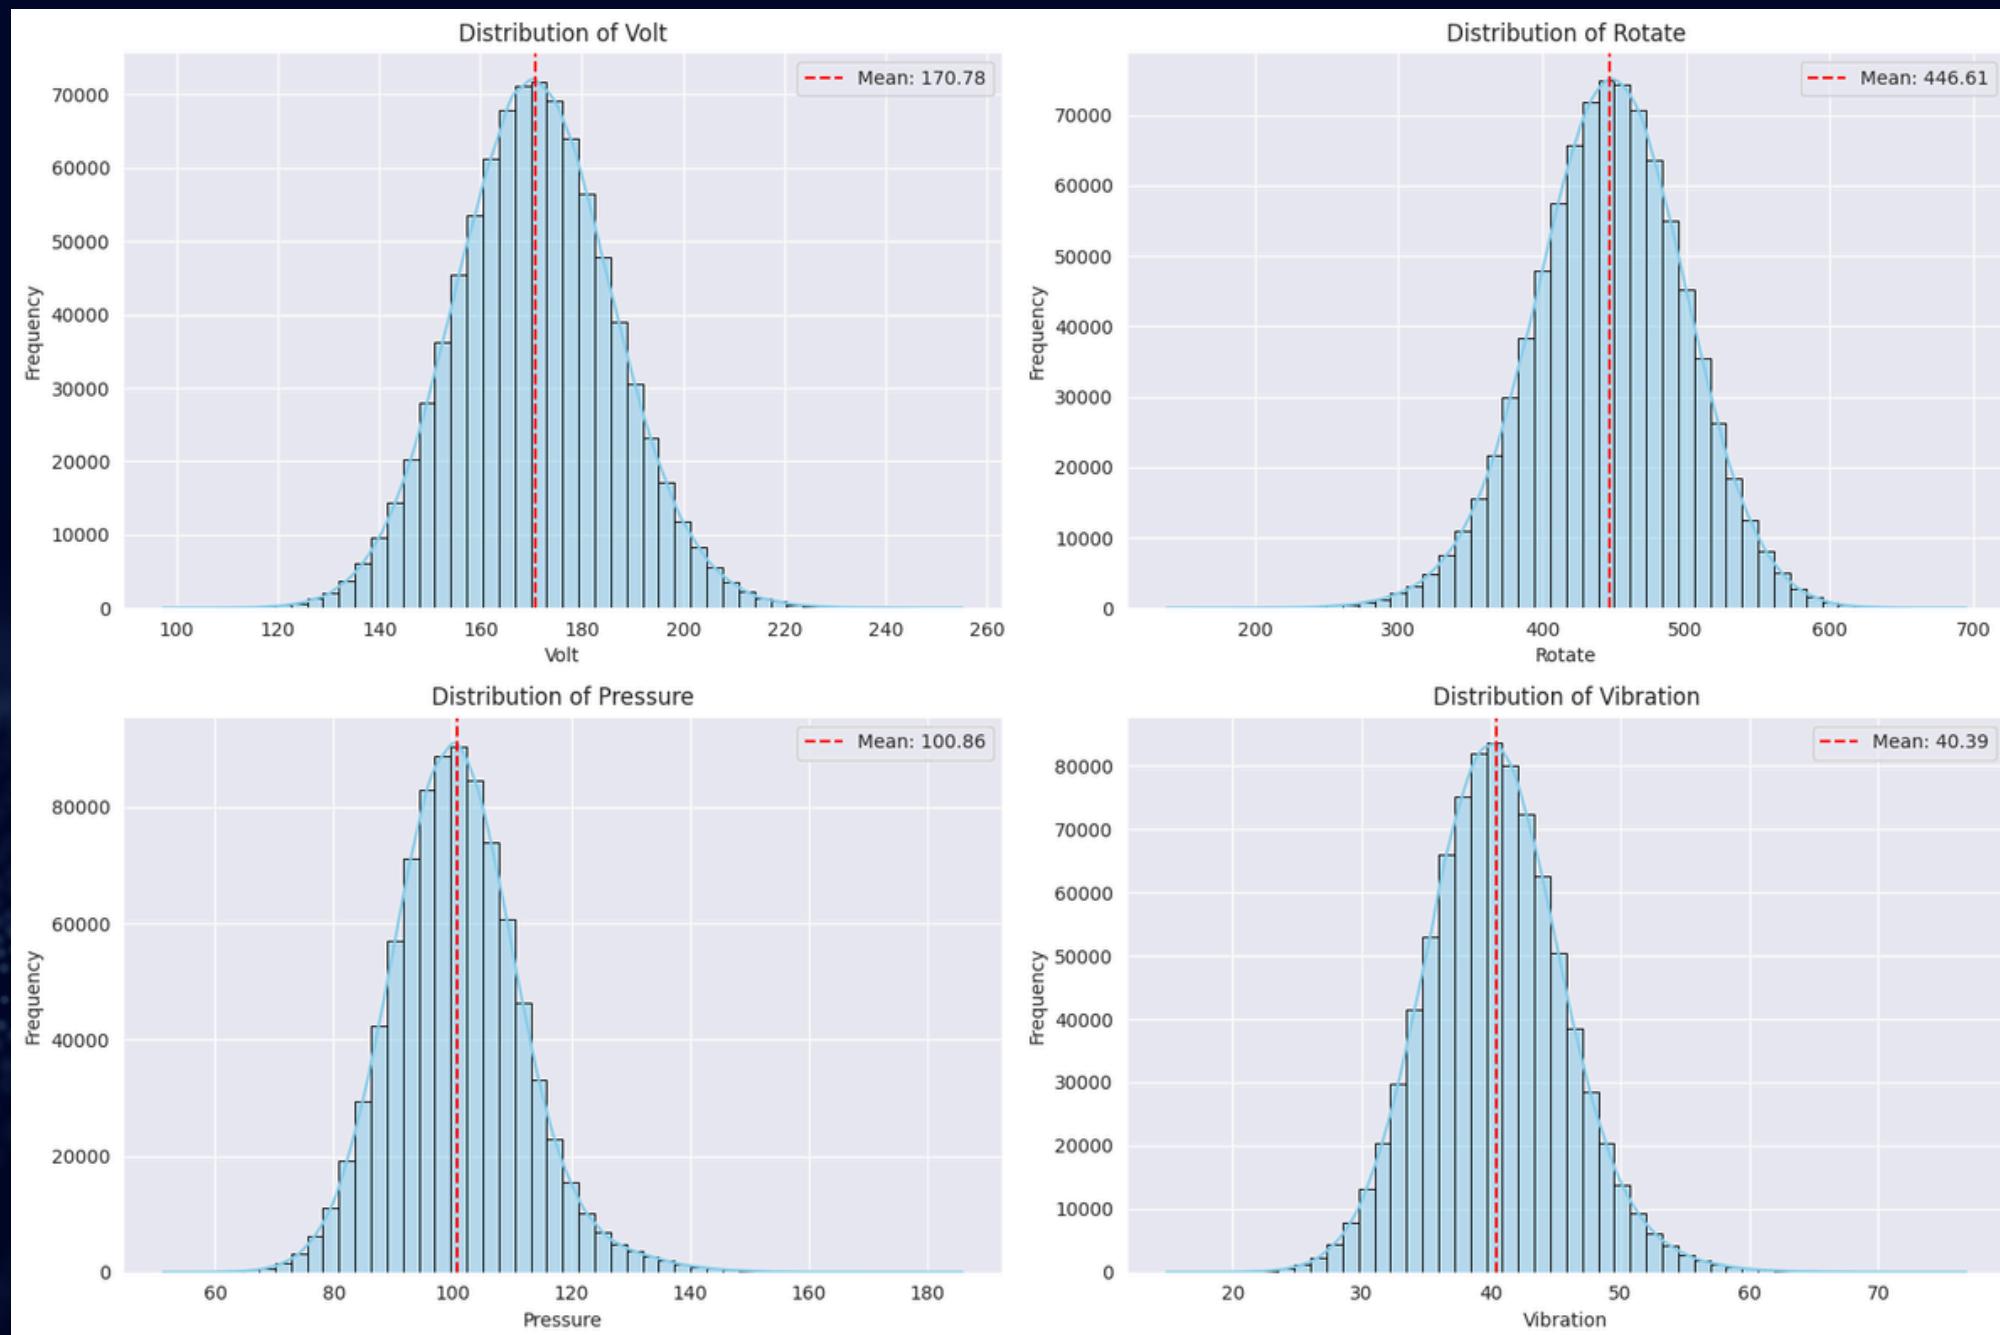


Fig. 2: Sensor Outliers Visualization (telemetry)



==== Skewness & Kurtosis ====
 volt: Skewness = 0.09, Kurtosis = 3.12
 rotate: Skewness = -0.14, Kurtosis = 3.20
 pressure: Skewness = 0.40, Kurtosis = 3.88
 vibration: Skewness = 0.25, Kurtosis = 3.47

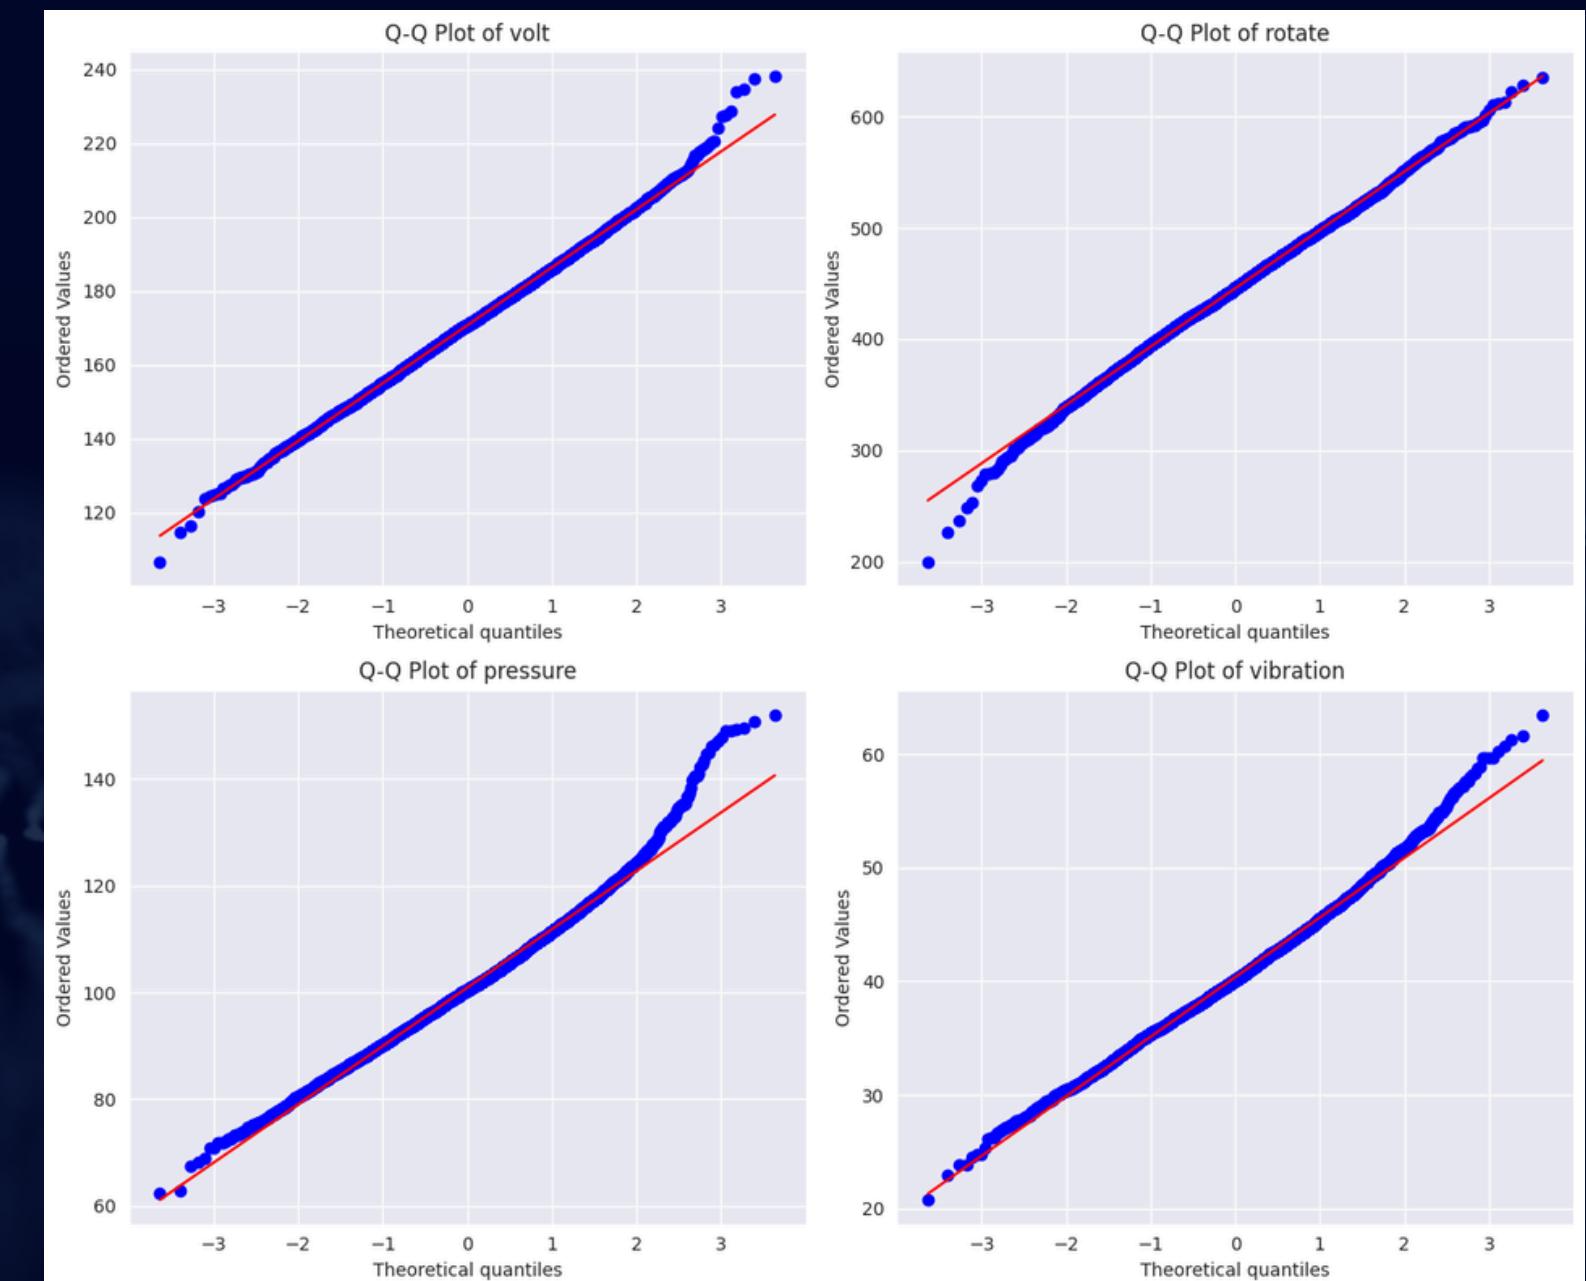


Fig. 3: Sensor Value Distributions with Mean Lines (telemetry)

Fig. 4: Skewness, Kurtosis & Q-Q Plots

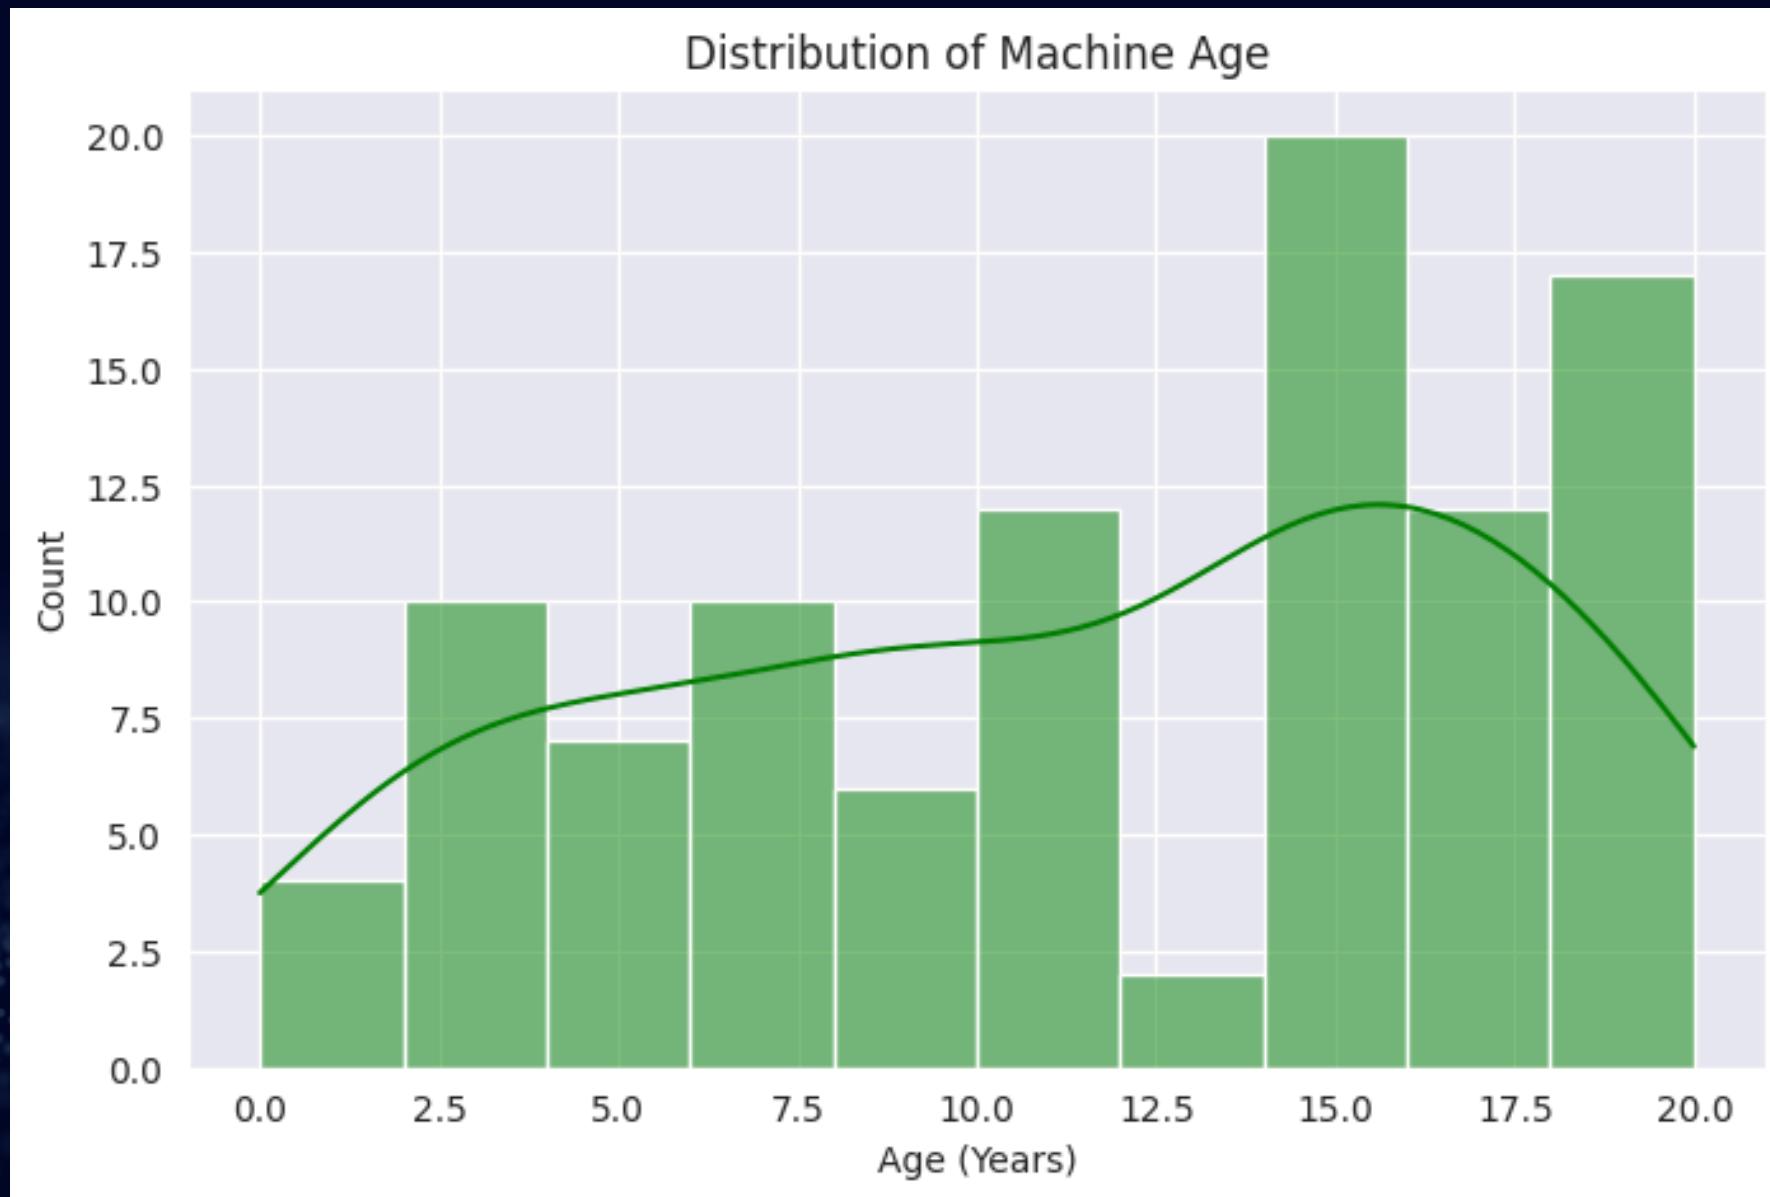


Fig. 5: Machine Age Distribution (machines)

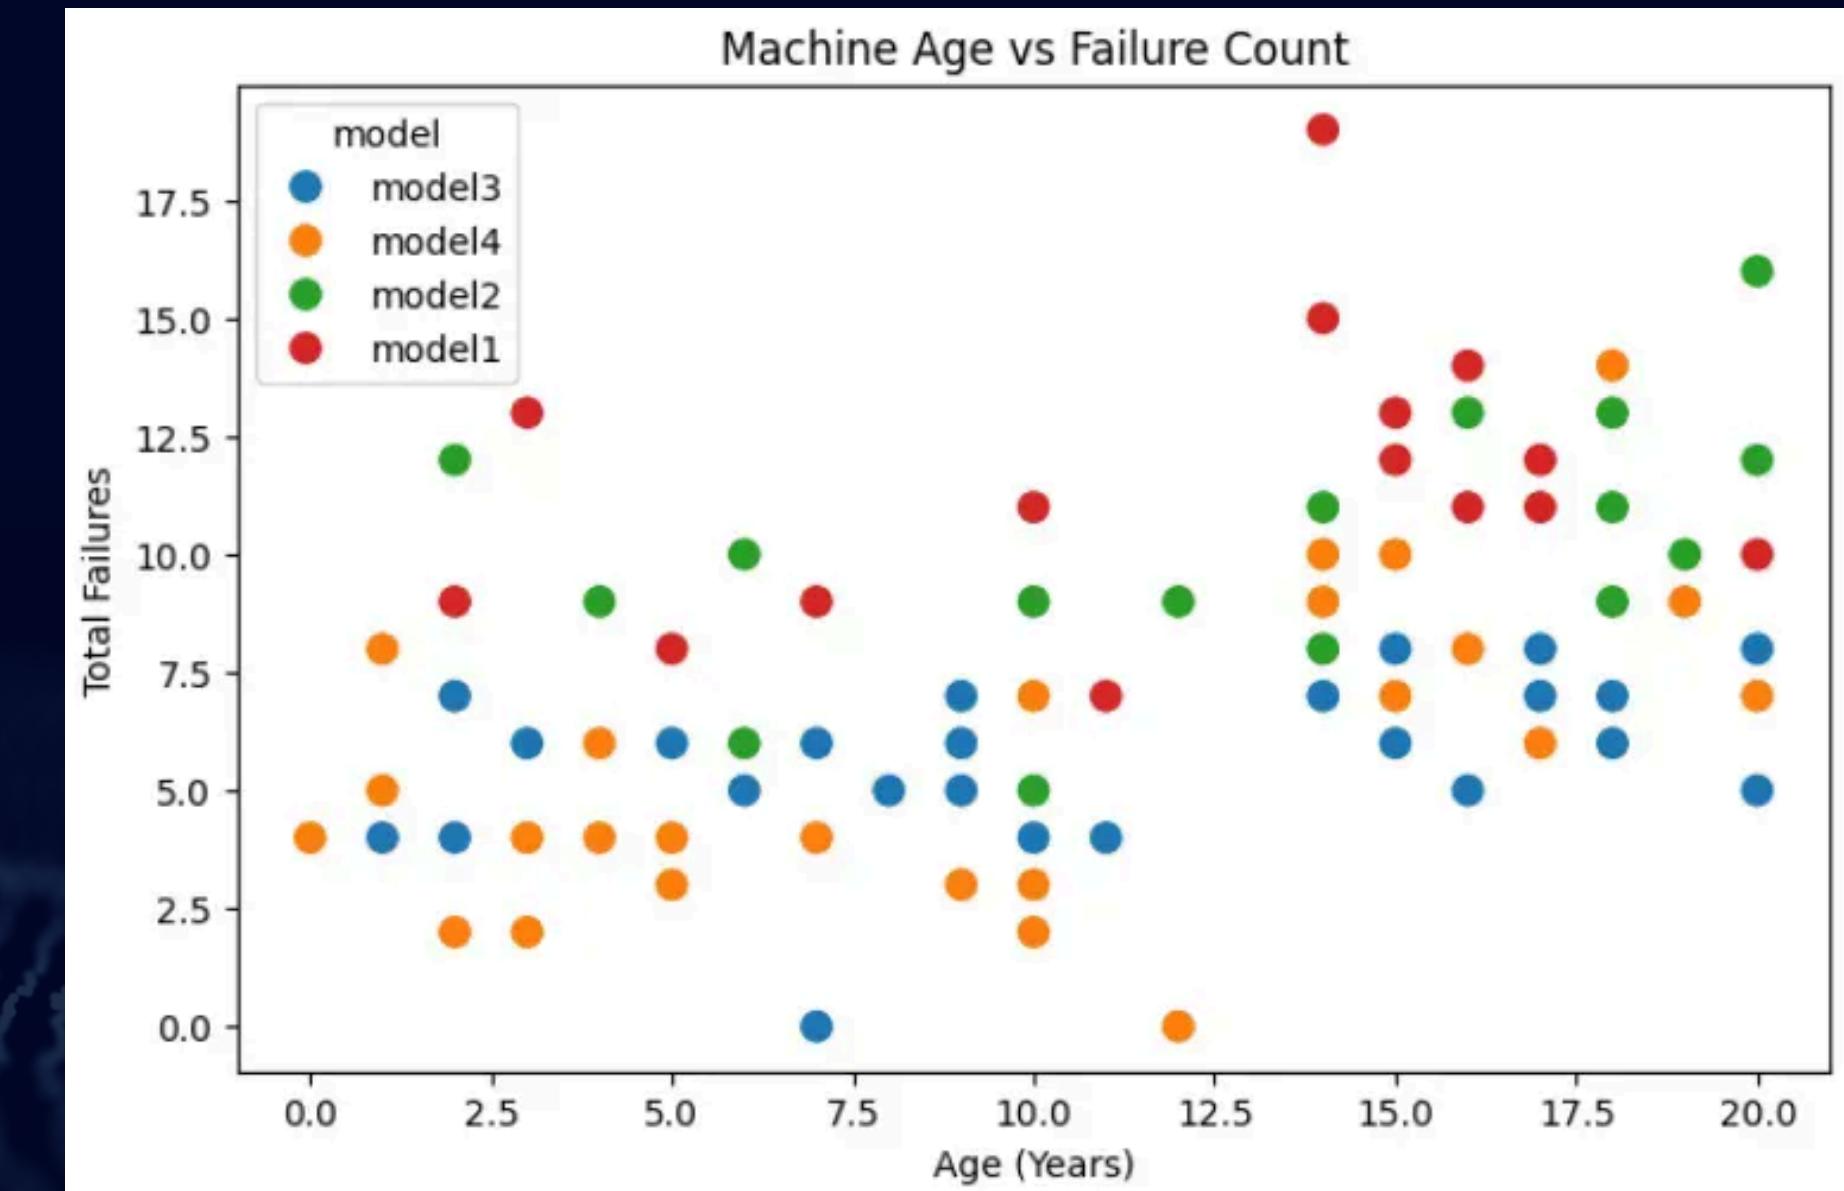


Fig. 6: Machine Age vs Failure Count

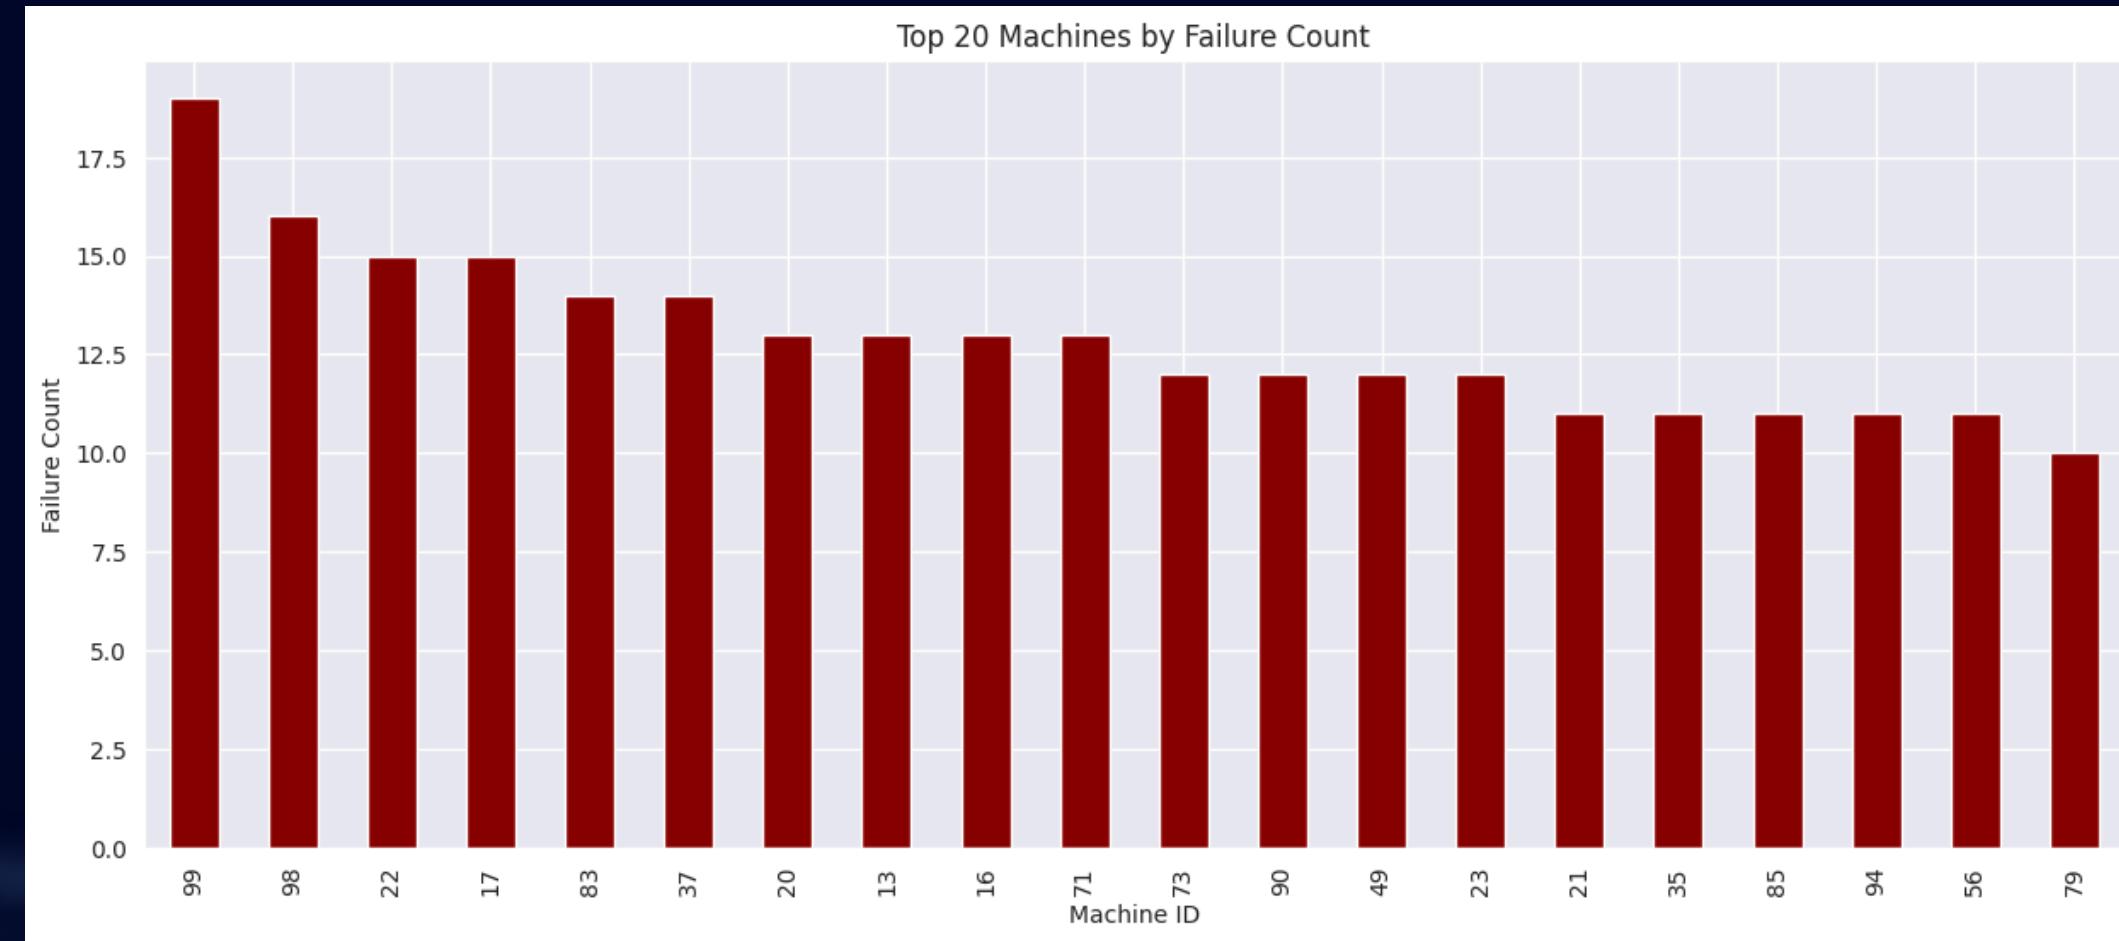


Fig 7: Top 20 Machines by Failrues

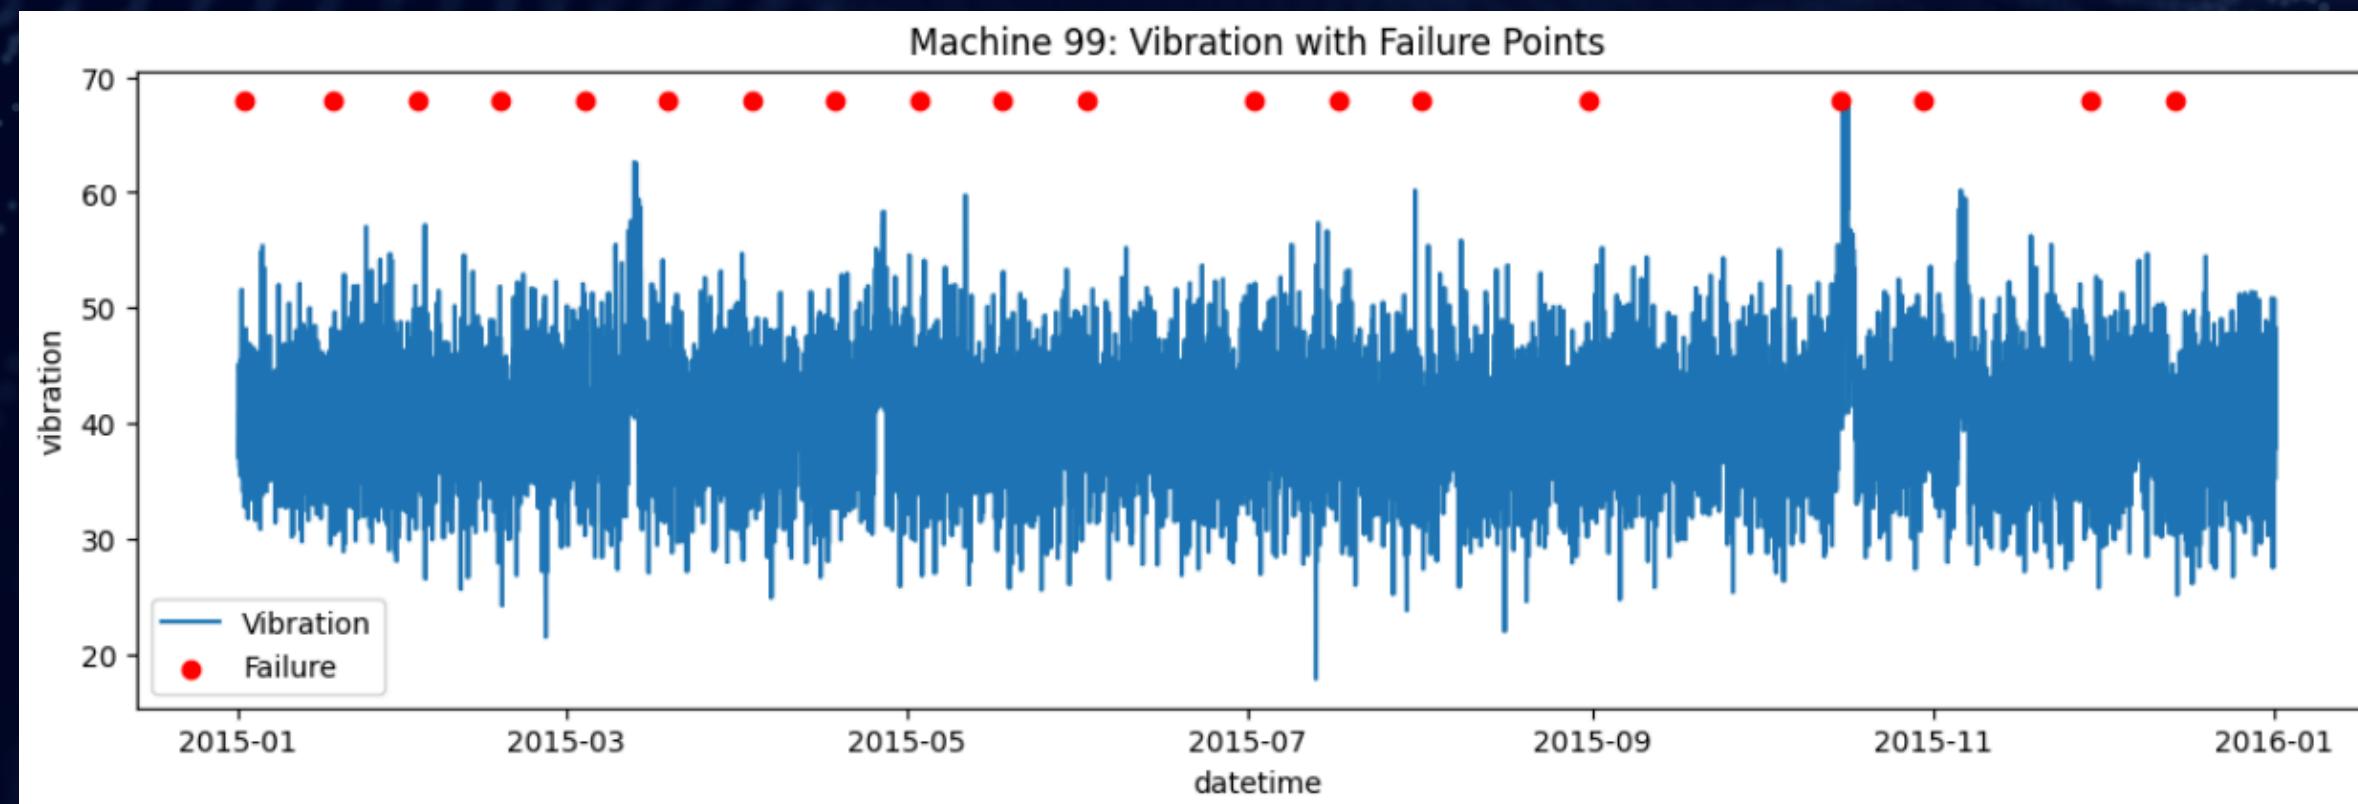


Fig 8: Vibration with Failure Points (machine 99)

Step 2: Data Preparation

1. Datasets Merge:

- Combine telemetry and machine metadata using machine ID
- Convert machine models into separate 0/1 columns (one-hot encoding)
- Convert error types into 0/1 columns
- Aggregate error data: Sum errors per machine and timestamp
- Combine all data into a single dataset

2. Target Variable Creation:

- Failure alignment: Matched telemetry rows with the next known failure for each machine
- Aggregated failures by machine and time
- Created binary labels: `will_fail_24h`, `will_fail_48h`, `will_fail_72h`
- Final dataset: Combined sensor readings, machine info, error data, and failure labels

```
will_fail_24h    17183
will_fail_48h    33959
will_fail_72h    50735
dtype: int64
will_fail_24h  will_fail_48h  will_fail_72h
0              0              0          825365
1              1              1          17183
0              0              1          16776
                  1              1          16776
Name: count, dtype: int64
Percentage of positive cases: 1.96%
Percentage of positive cases: 3.88%
Percentage of positive cases: 5.79%
```

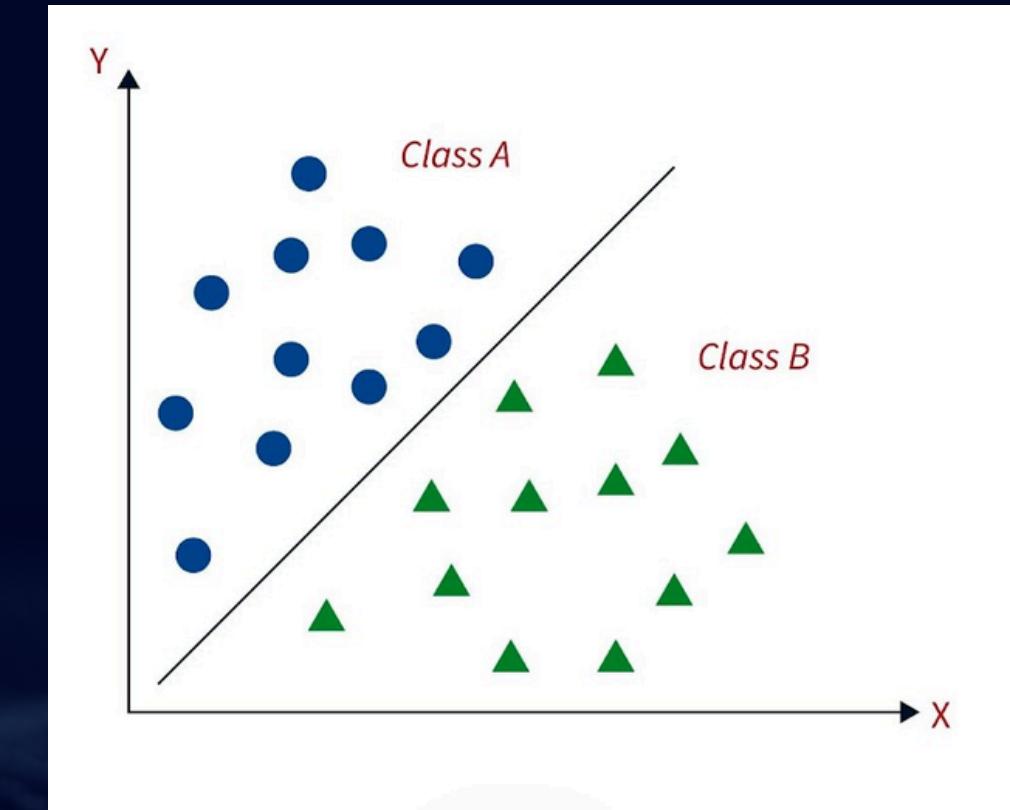


Fig 10: Distribution of Positive vs Negative Cases for Failure Prediction

	datetime	machineID	volt	rotate	pressure	vibration	age	model1	model2	model3	model4	error1	error2	error3	error4	error5	will_fail_24h	will_fail_48h	will_fail_72h
0	2015-01-01 06:00:00	1	176.217853	418.504078	113.077935	45.087686	18	0	0	1	0	0	0	0	0	0	0	0	
1	2015-01-01 07:00:00	1	162.879223	402.747490	95.460525	43.413973	18	0	0	1	0	0	0	0	0	0	0	0	
2	2015-01-01 08:00:00	1	170.989902	527.349825	75.237905	34.178847	18	0	0	1	0	0	0	0	0	0	0	0	
3	2015-01-01 09:00:00	1	162.462833	346.149335	109.248561	41.122144	18	0	0	1	0	0	0	0	0	0	0	0	
4	2015-01-01 10:00:00	1	157.610021	435.376873	111.886648	25.990511	18	0	0	1	0	0	0	0	0	0	0	0	

```
Missing values per column:
datetime          0
machineID         0
volt              0
rotate             0
pressure           0
vibration          0
age                0
model1             0
model2             0
model3             0
model4             0
error1             0
error2             0
error3             0
error4             0
error5             0
will_fail_24h      0
will_fail_48h      0
will_fail_72h      0
dtype: int64
```

Number of duplicate rows: 0

Checked dataset for missing values → none found

Checked for duplicate rows → none found

Ensured dataset is clean and ready for modeling

Step 3: Feature Engineering

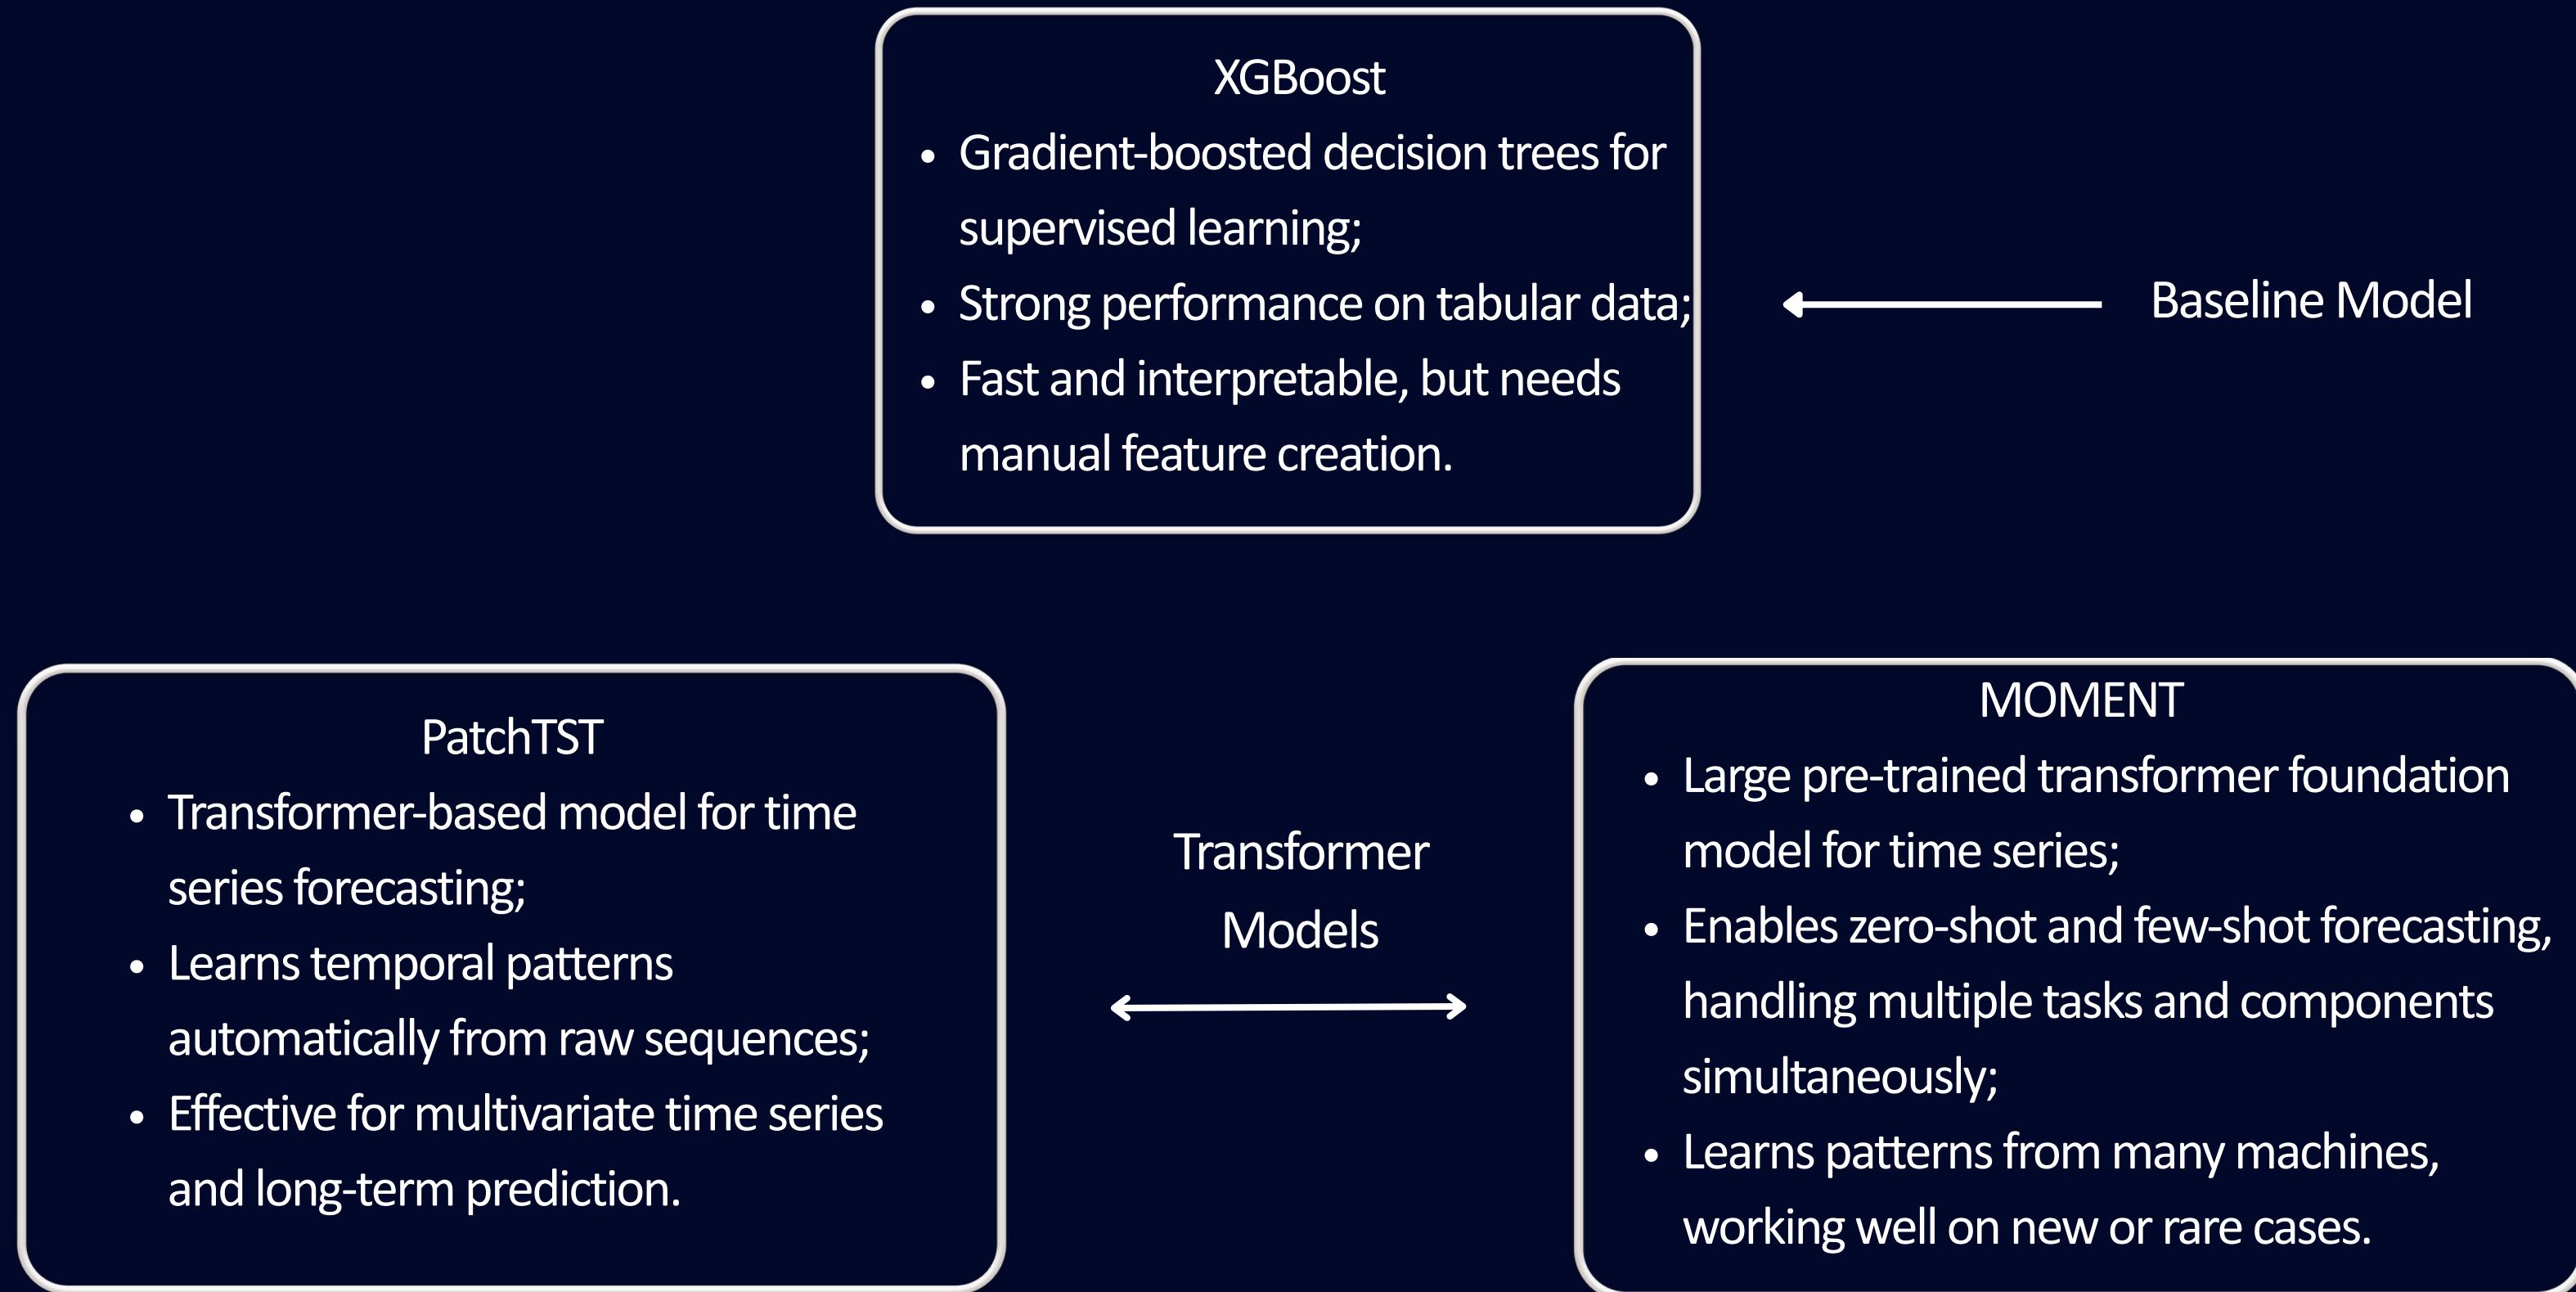
- Focused on sensor data: voltage, rotation, pressure, vibration
- Computed rolling window features (6h, 12h, 24h)
- Mean, Standard Deviation, Min, Max
- Captures short-term and long-term trends in machine behavior
- Missing values handled with forward fill → 0
- Reason: Needed additional features for XGBoost model, which relies on manual feature input

Step 4: Data Splitting & Model Preparation

- Sorted data by machine ID and datetime
- Time-based split:
 - Train: first 70% of data
 - Validation: next 15% of data
 - Test: final 15% of data
- Checked class imbalance in each split for 24h, 48h, 72h prediction windows

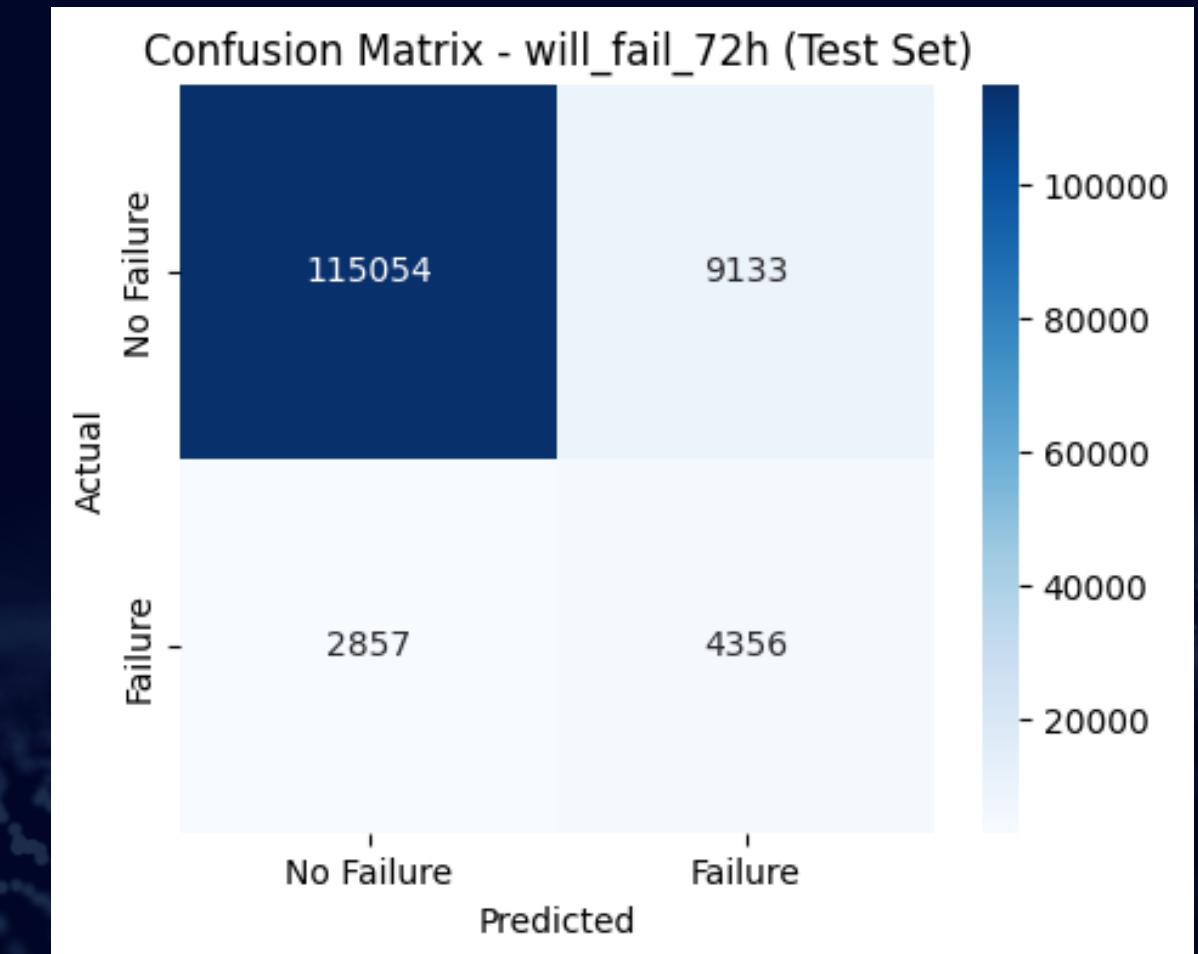
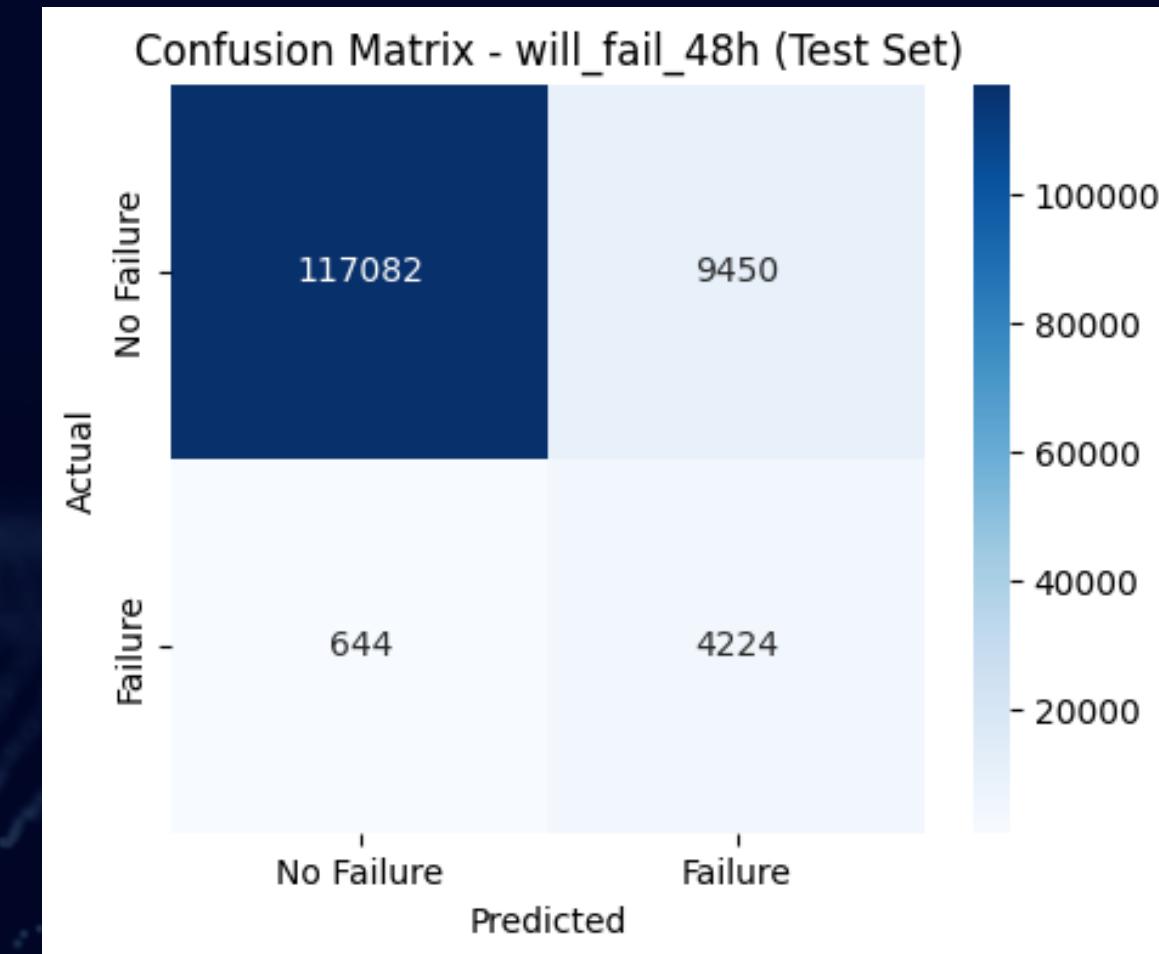
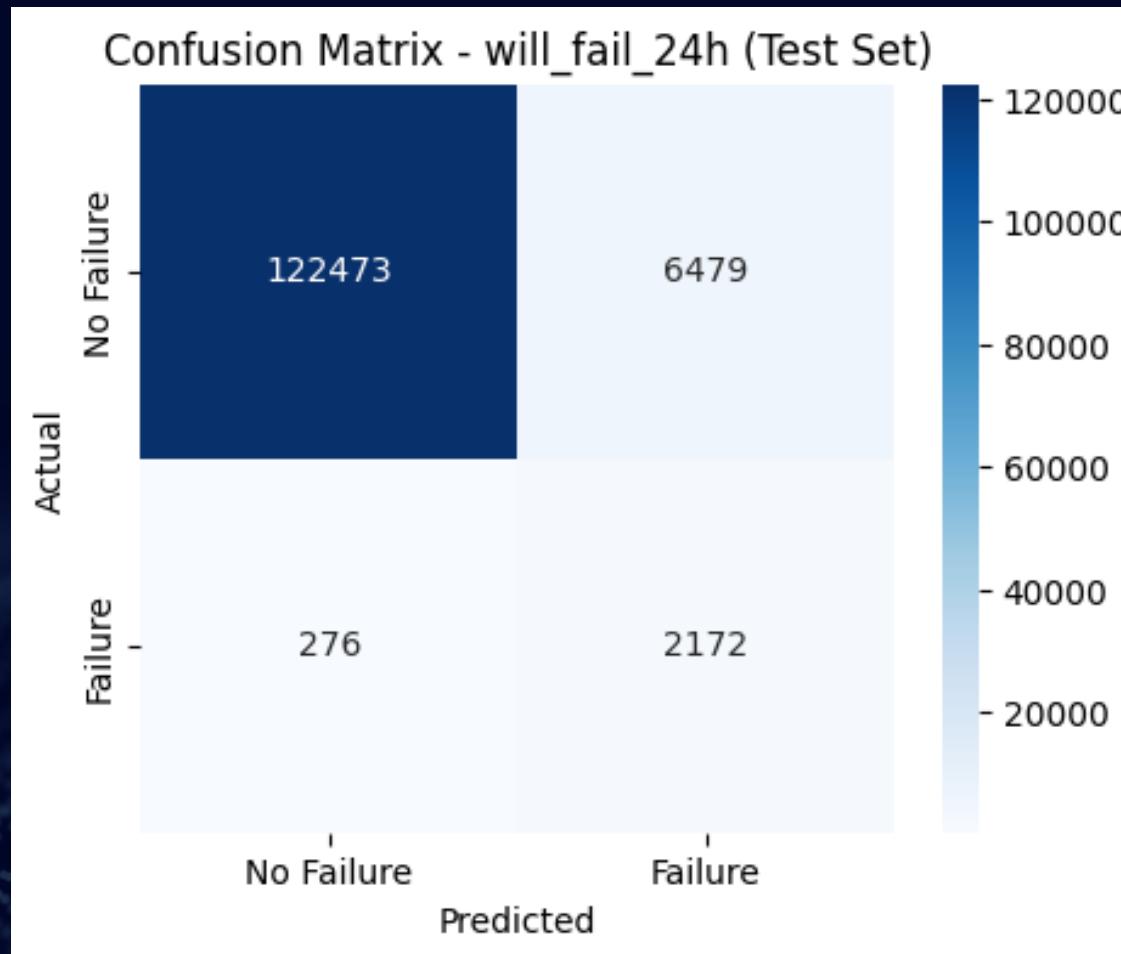
```
train: n=613,300  time 2015-01-01 06:00:00 → 2015-09-13 18:00:00
      will_fail_24h: positives=12,542  rate=2.045%
      will_fail_48h: positives=24,723  rate=4.031%
      will_fail_72h: positives=36,928  rate=6.021%
val: n=131,400   time 2015-09-13 19:00:00 → 2015-11-07 12:00:00
      will_fail_24h: positives=2,193  rate=1.669%
      will_fail_48h: positives=4,368  rate=3.324%
      will_fail_72h: positives=6,594  rate=5.018%
test: n=131,400  time 2015-11-07 13:00:00 → 2016-01-01 06:00:00
      will_fail_24h: positives=2,448  rate=1.863%
      will_fail_48h: positives=4,868  rate=3.705%
      will_fail_72h: positives=7,213  rate=5.489%
```

Step 5: Model Selection



Results & Analysis

Model: XGBoost



== TEST SET REPORT ==

	precision	recall	f1-score	support
No Failure	1.00	0.95	0.97	128952
Failure	0.25	0.89	0.39	2448
accuracy			0.95	131400
macro avg	0.62	0.92	0.68	131400
weighted avg	0.98	0.95	0.96	131400

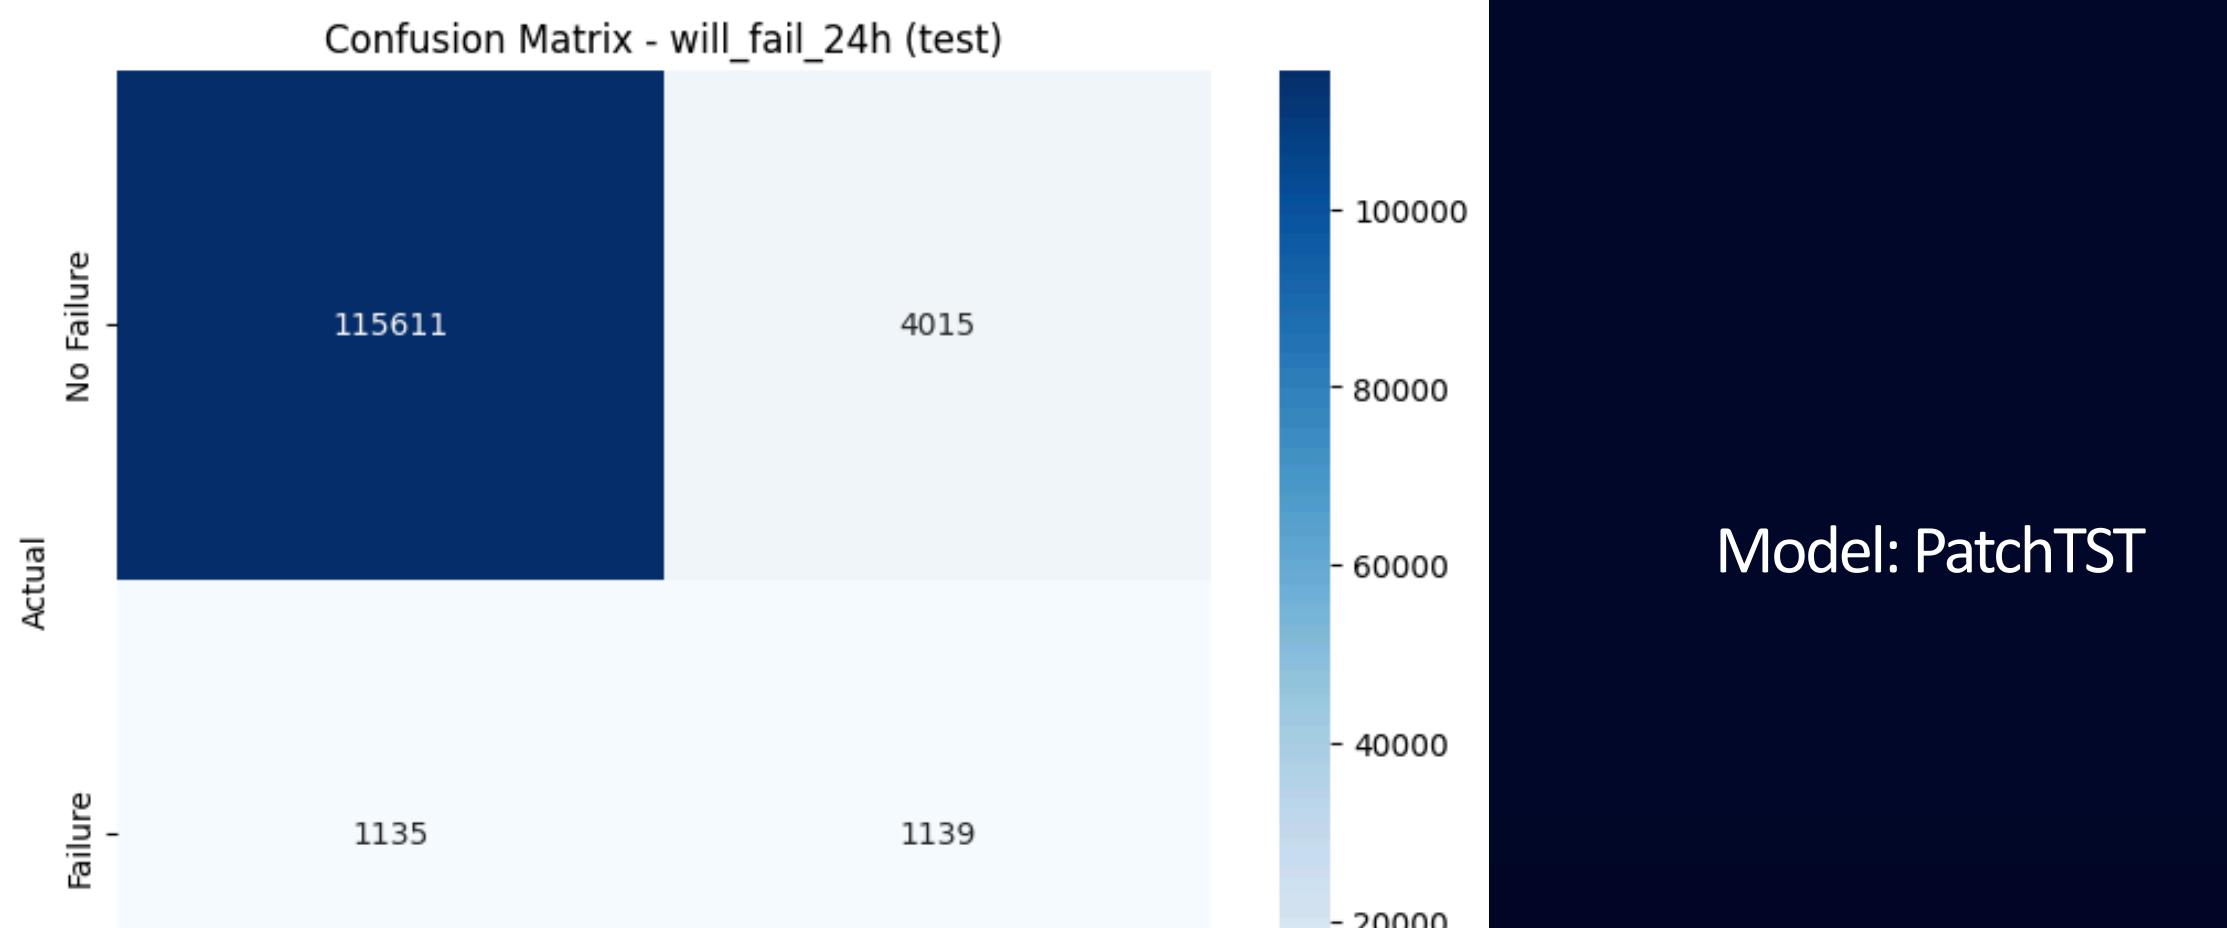
== TEST SET REPORT ==

	precision	recall	f1-score	support
No Failure	0.99	0.93	0.96	126532
Failure	0.31	0.87	0.46	4868
accuracy			0.92	131400
macro avg	0.65	0.90	0.71	131400
weighted avg	0.97	0.92	0.94	131400

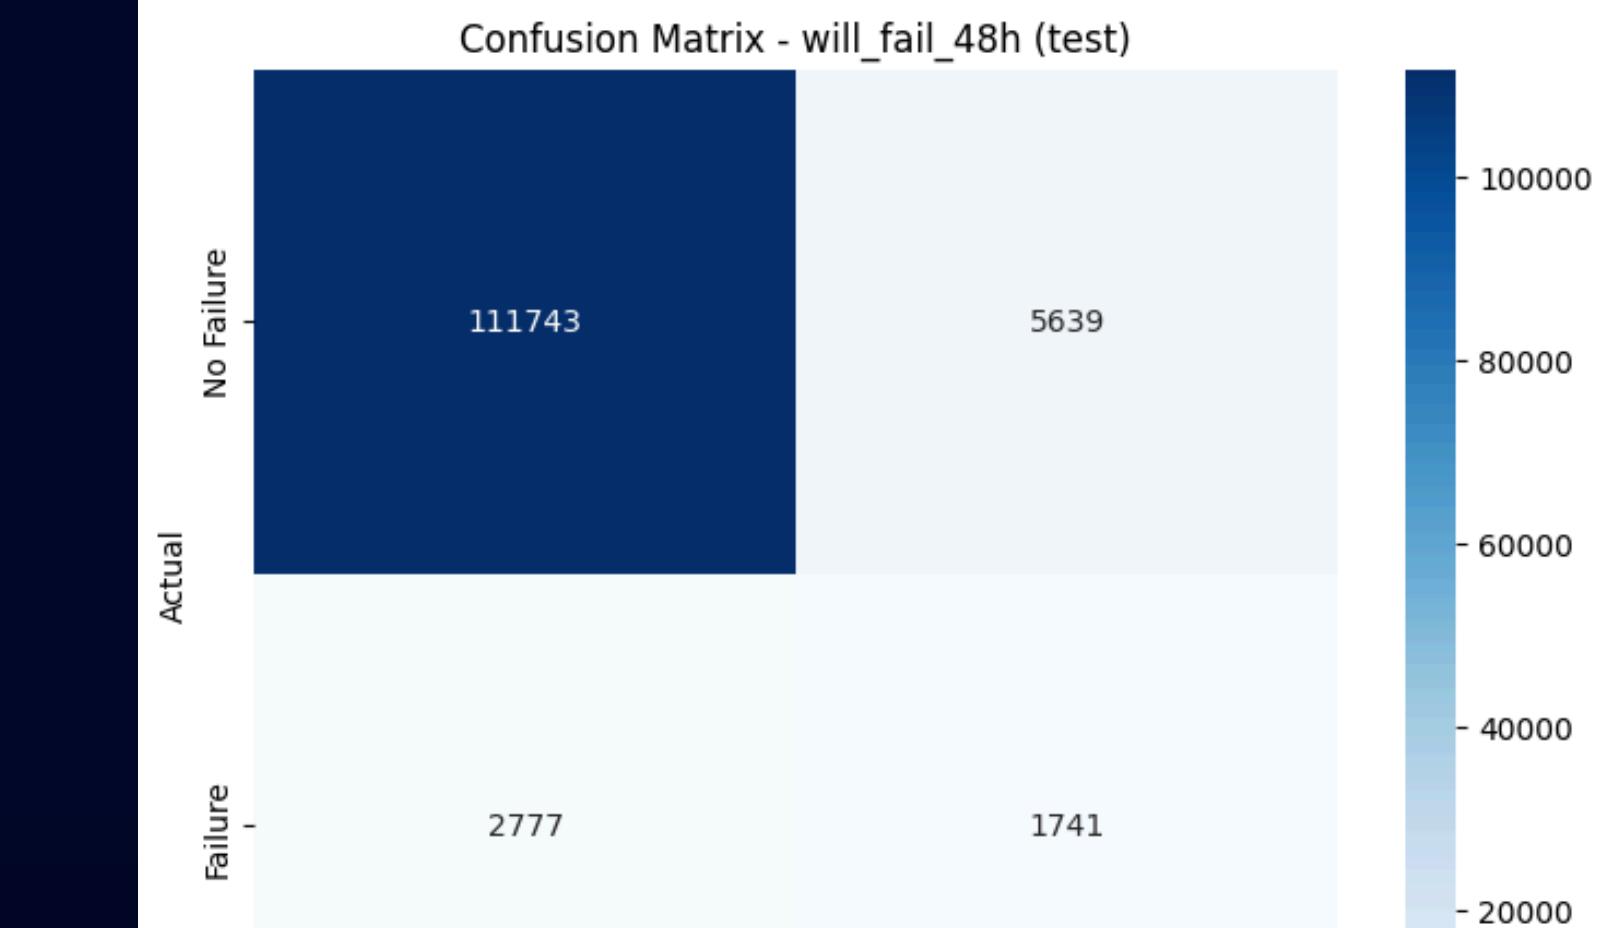
== TEST SET REPORT ==

	precision	recall	f1-score	support
No Failure	0.98	0.93	0.95	124187
Failure	0.32	0.60	0.42	7213
accuracy			0.91	131400
macro avg	0.65	0.77	0.69	131400
weighted avg	0.94	0.91	0.92	131400

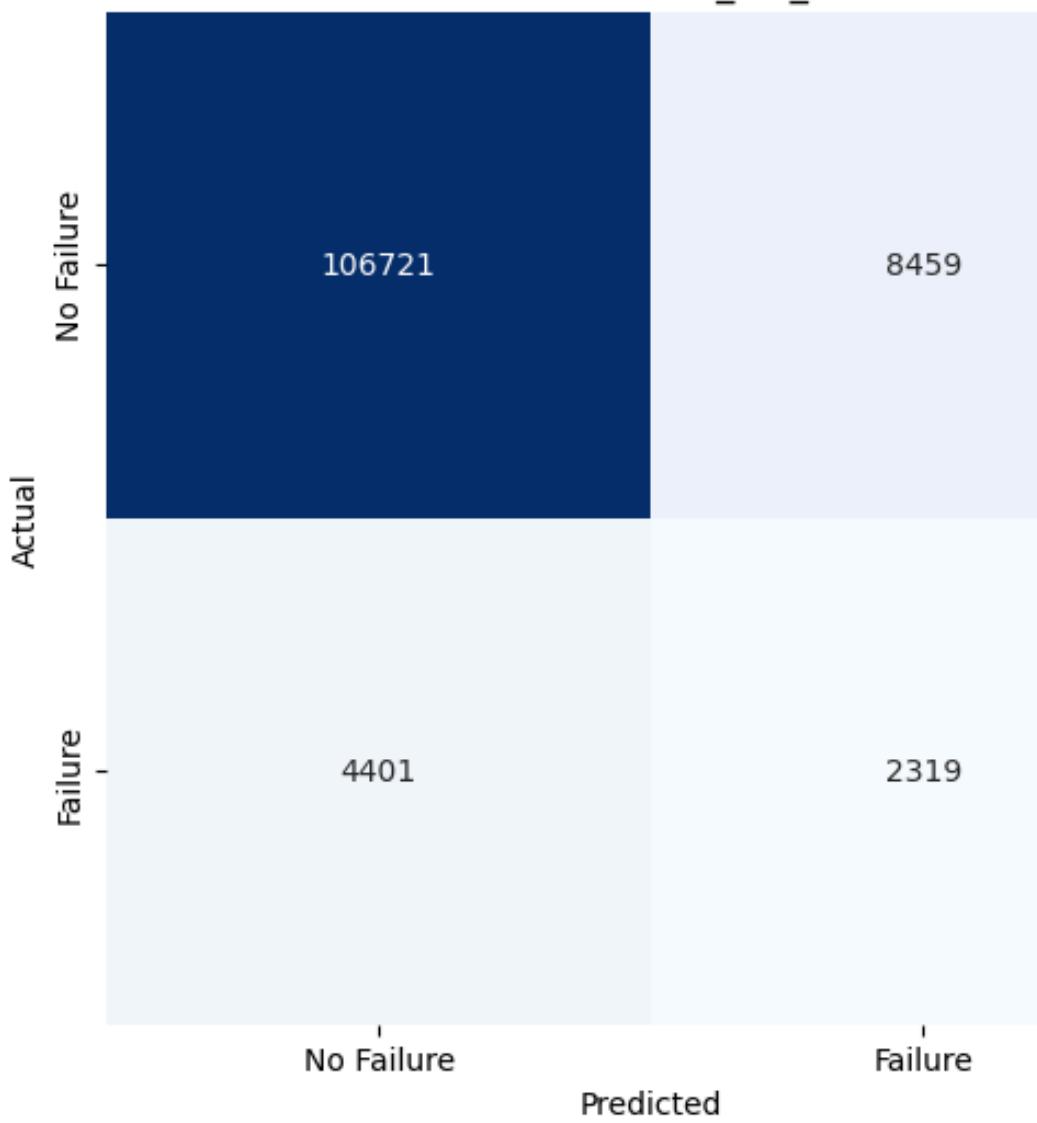
Confusion Matrix - will_fail_24h (test)



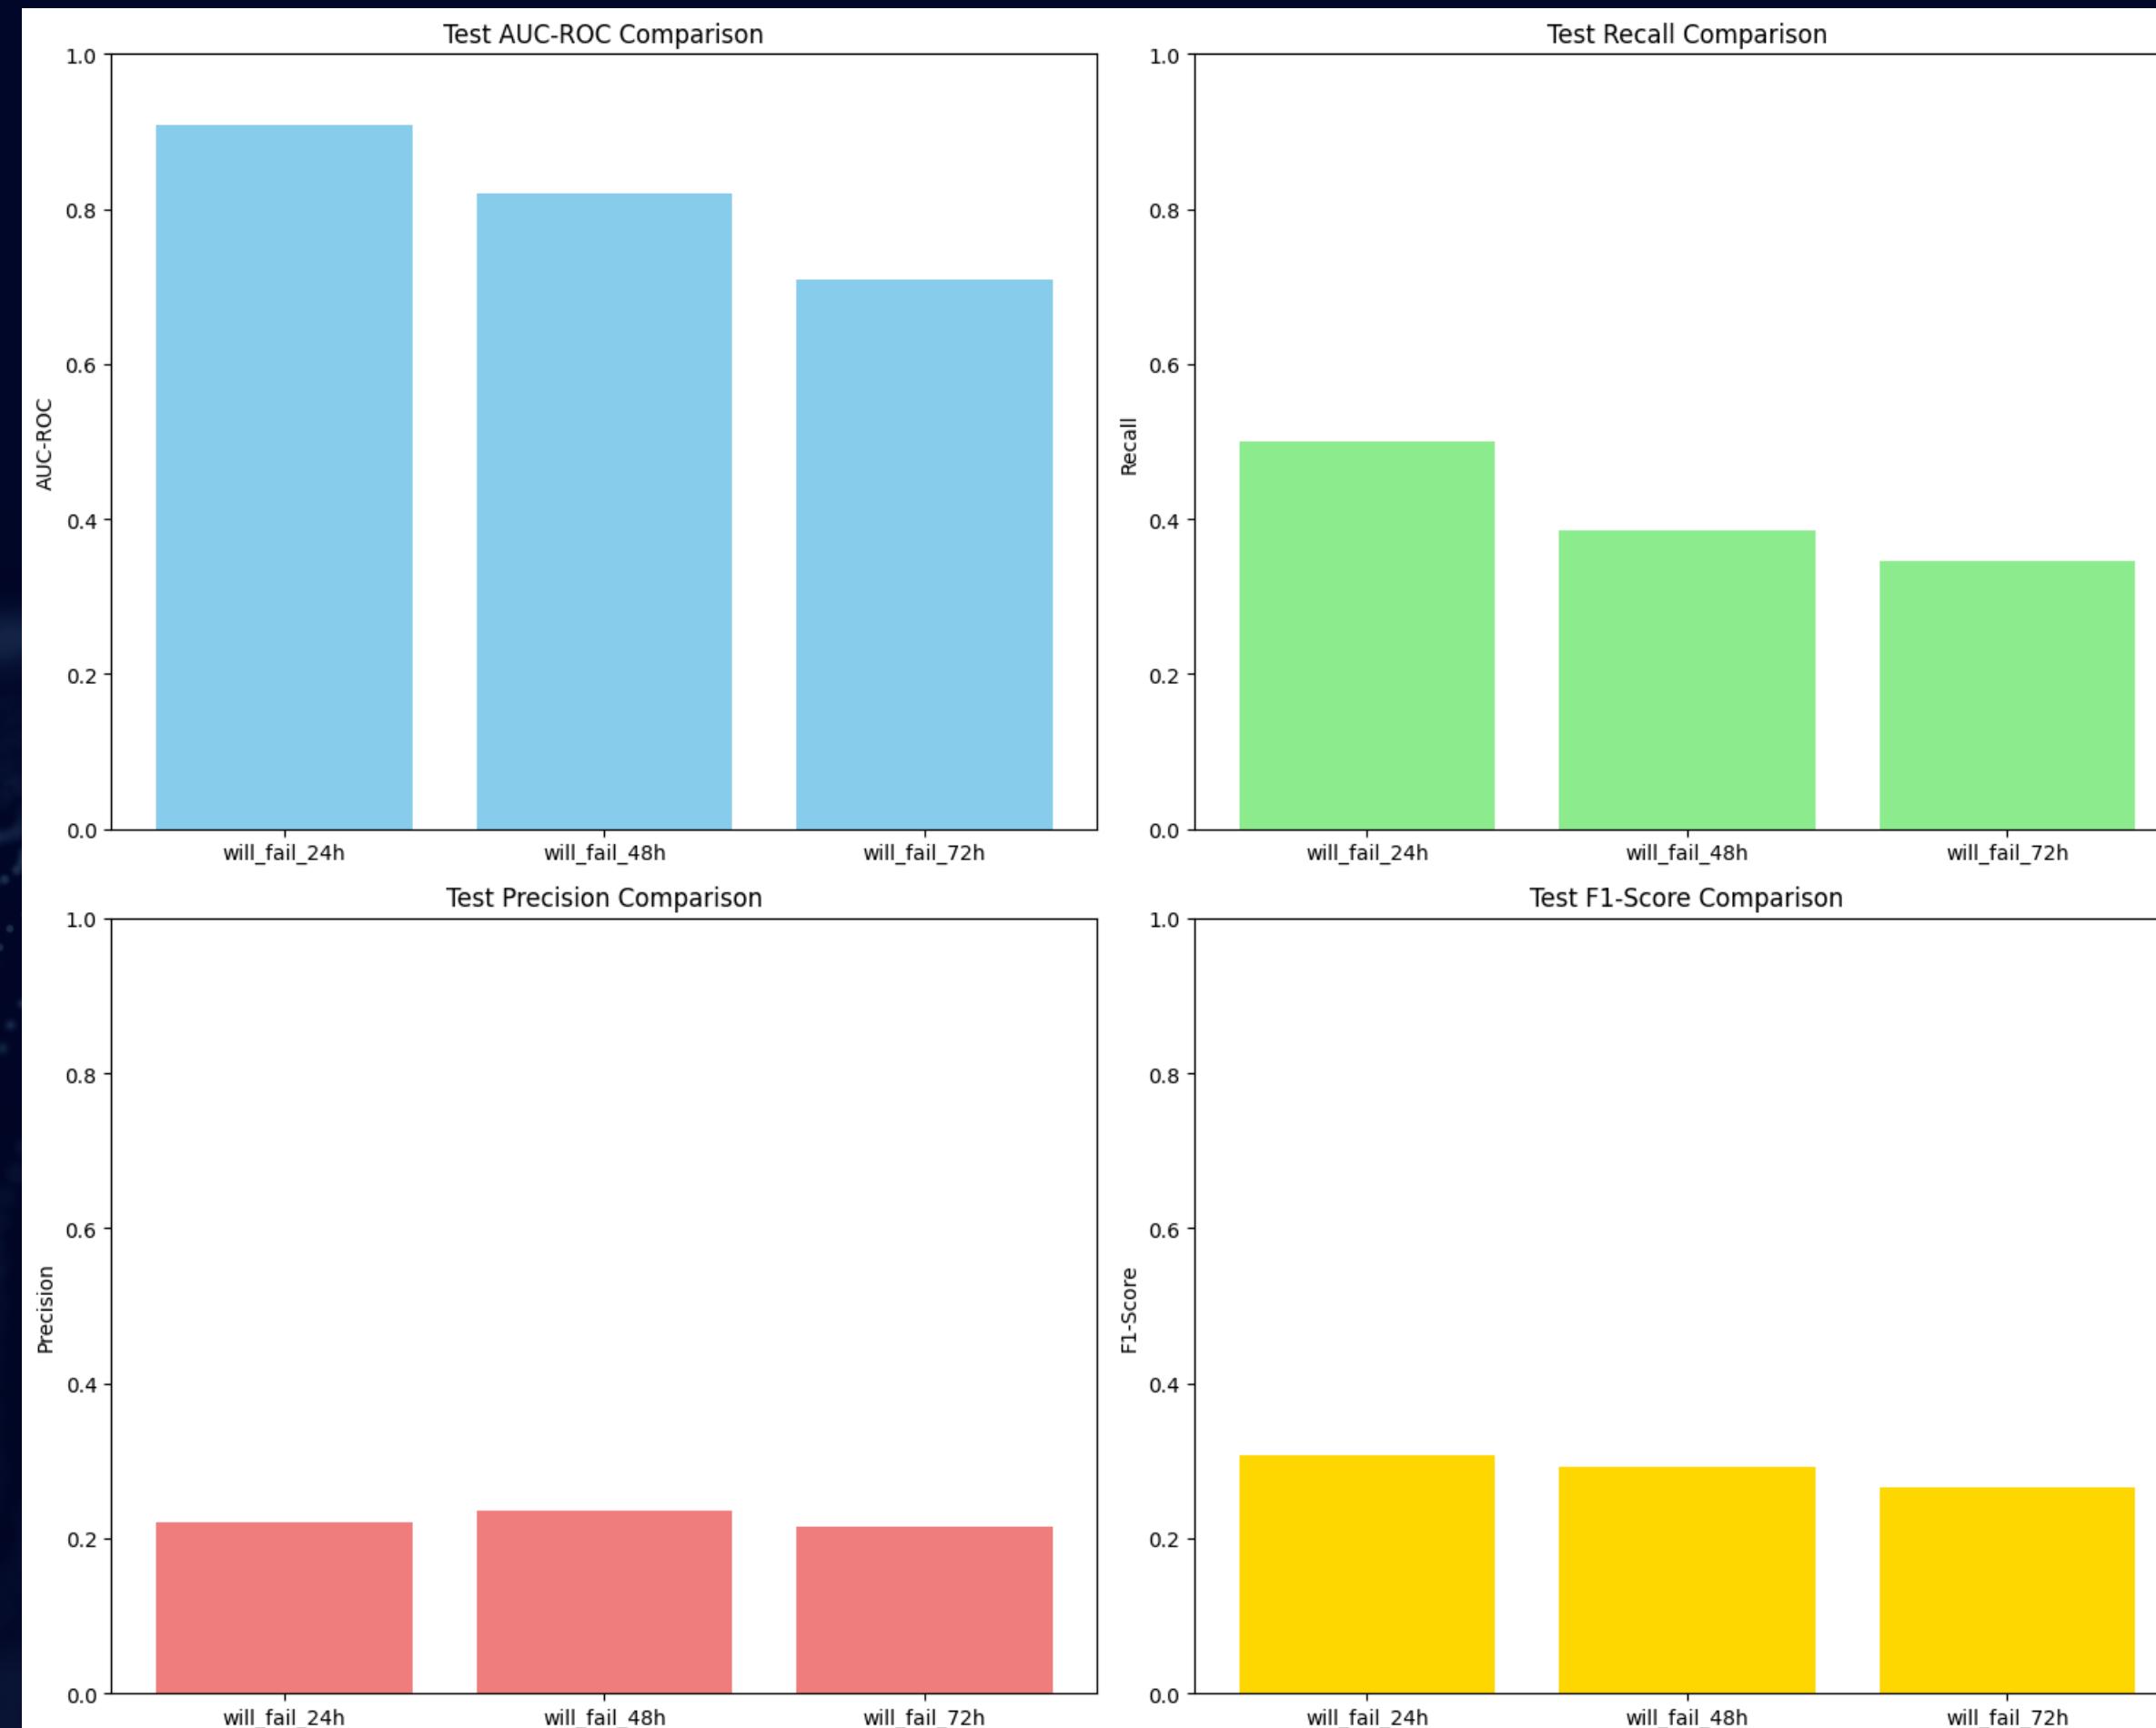
Confusion Matrix - will_fail_48h (test)



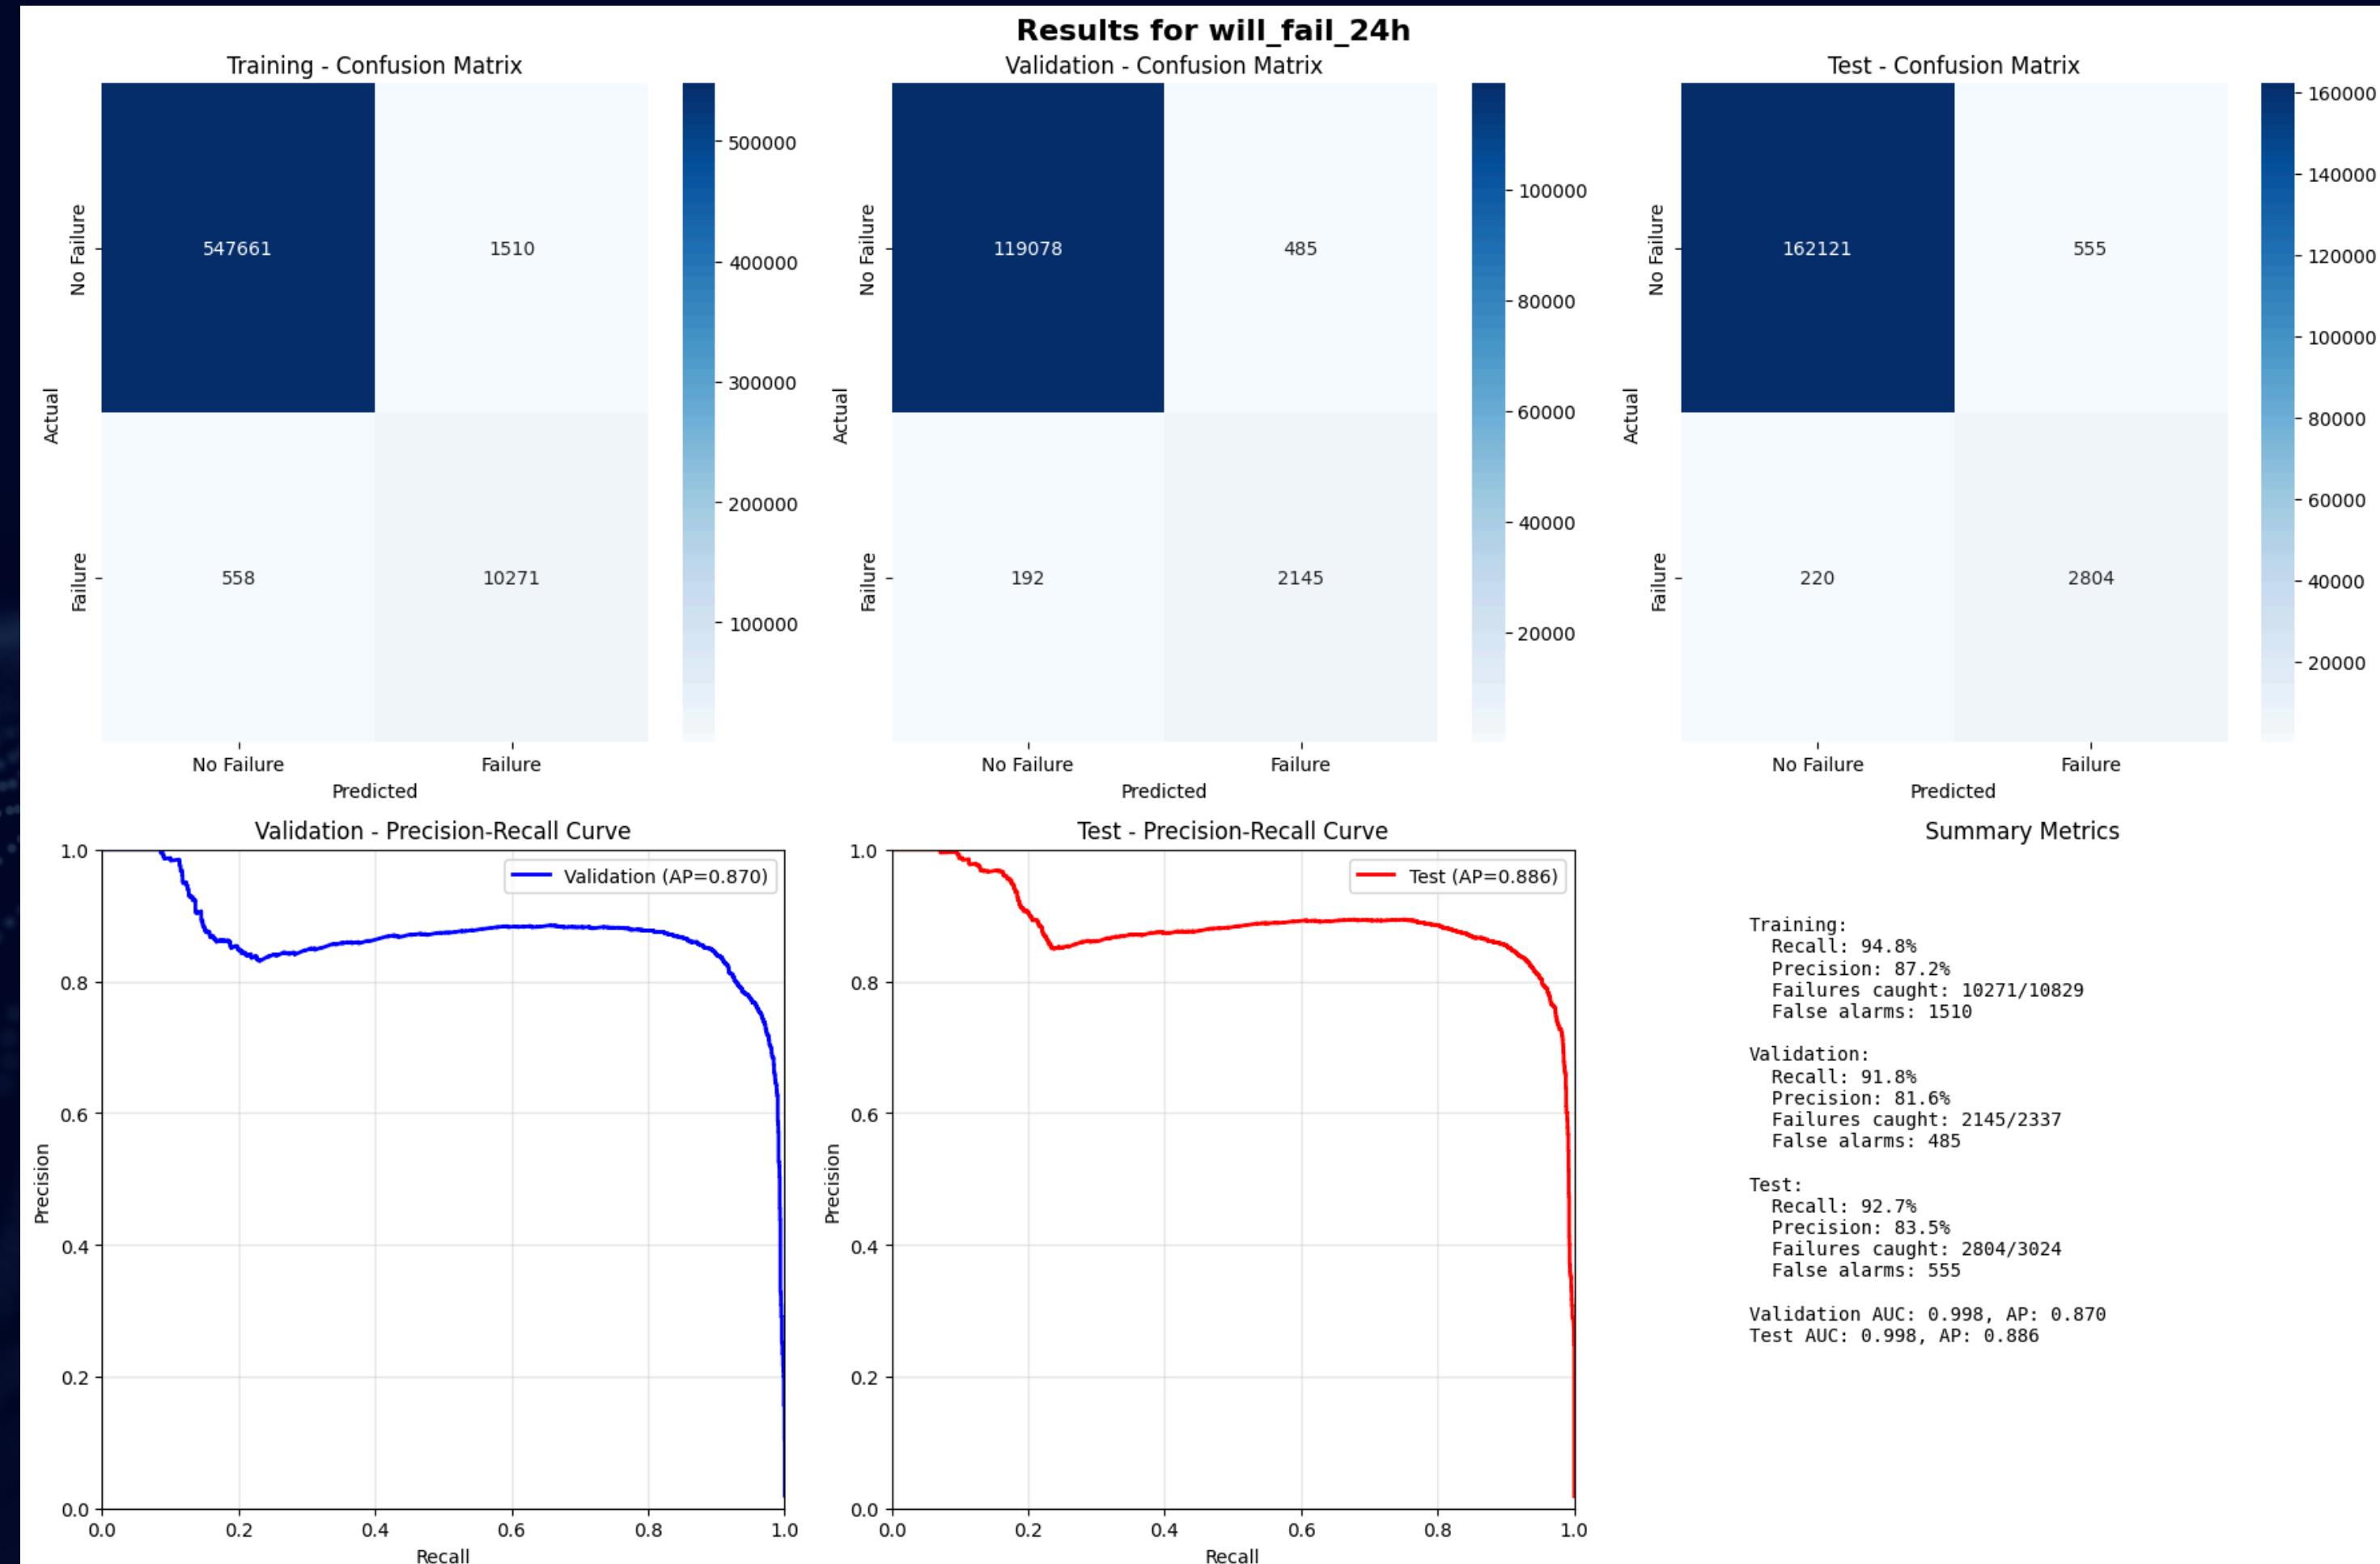
Confusion Matrix - will_fail_72h (test)



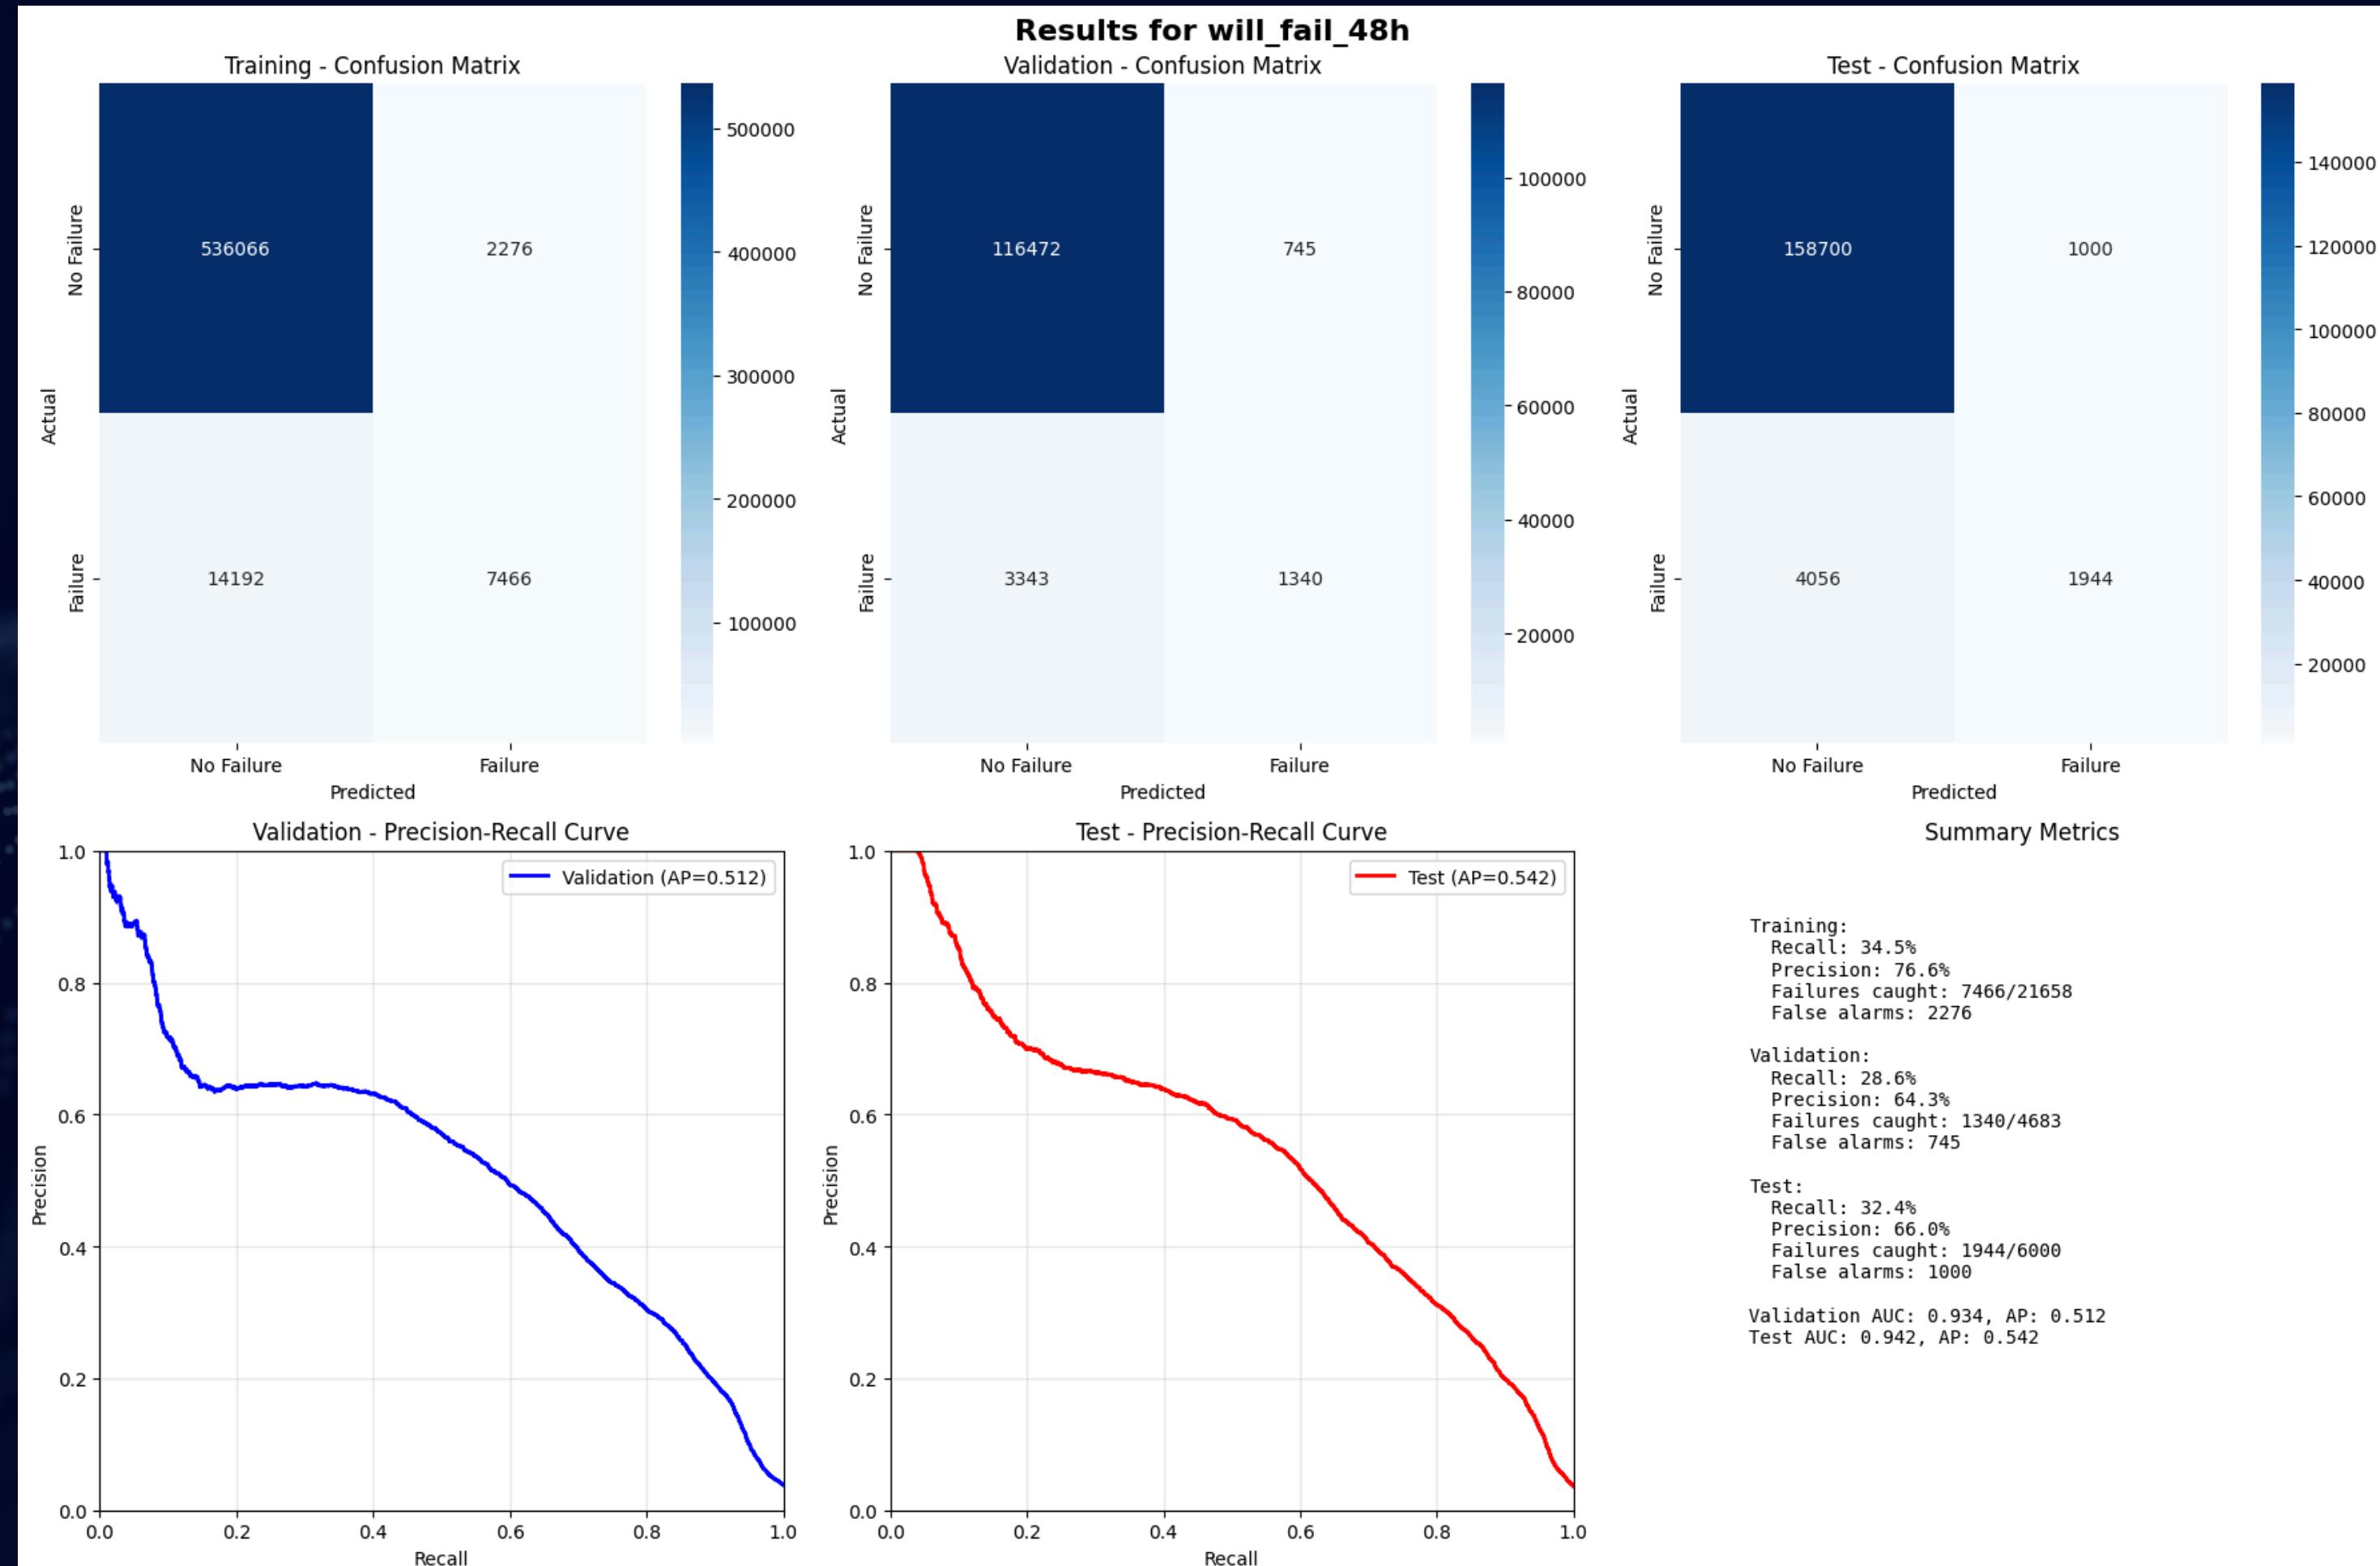
Model: PatchTST



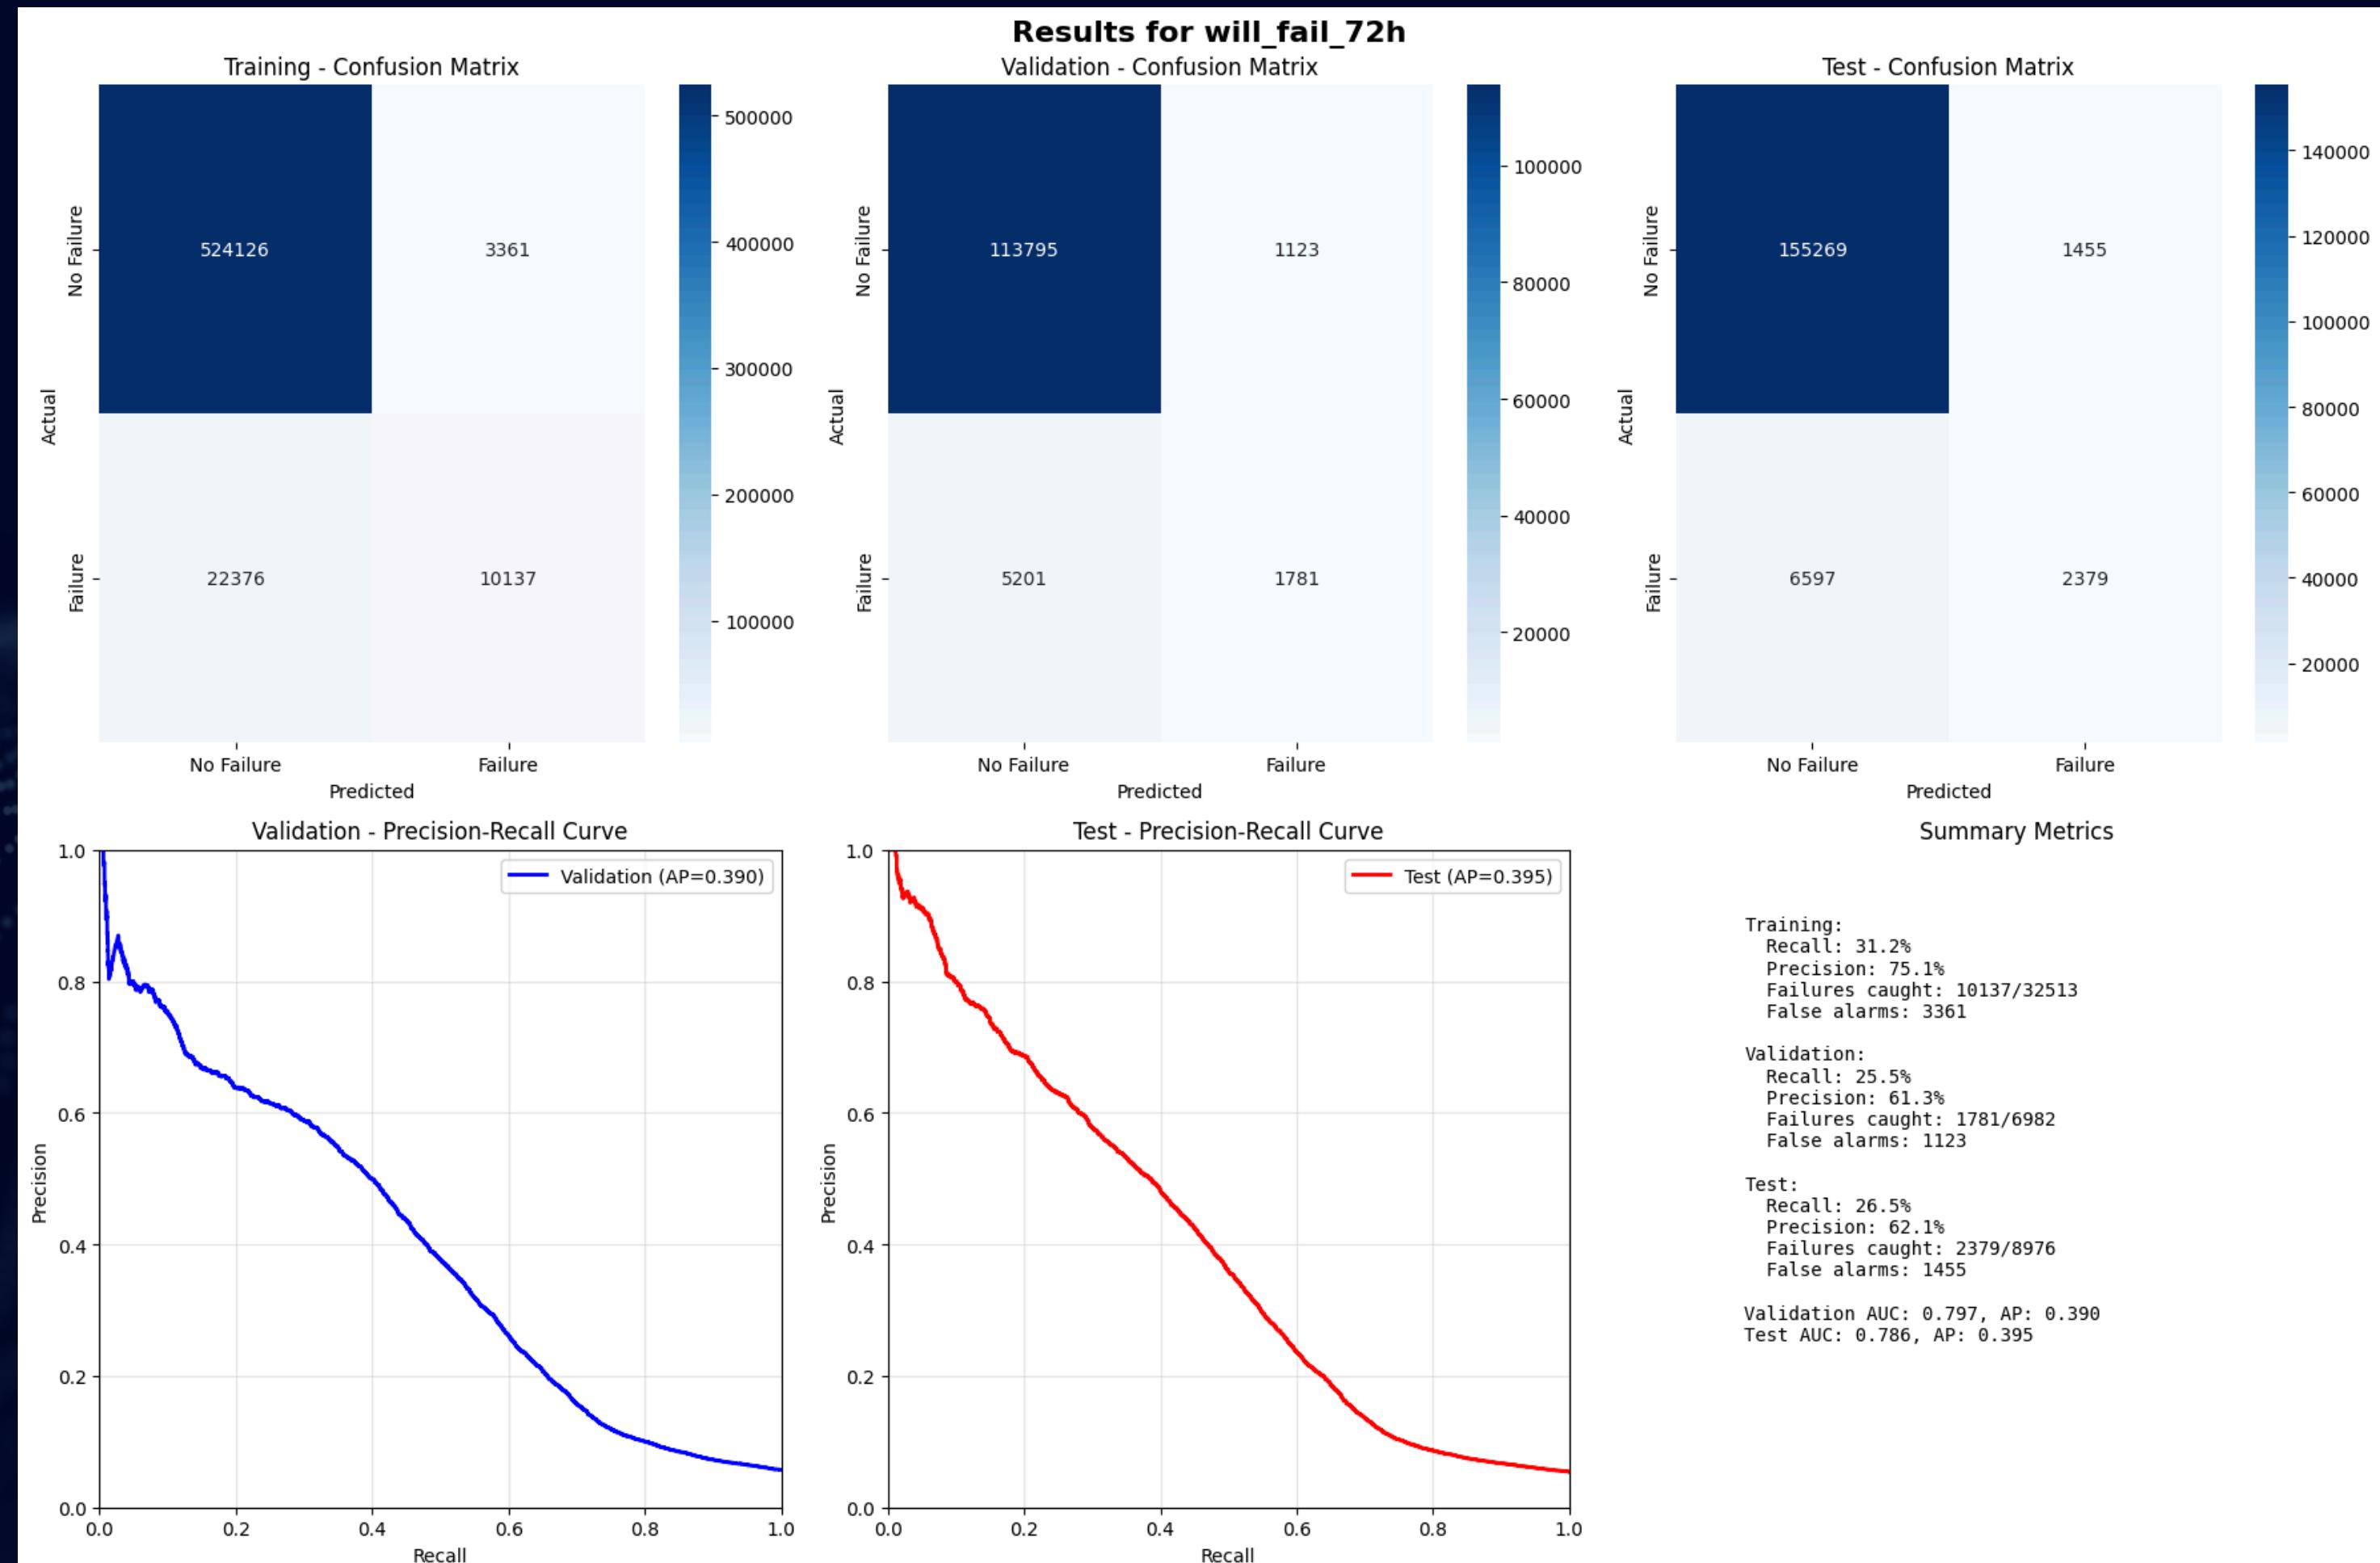
Model: PatchTST hybrid (24h)



Model: PatchTST hybrid (48h)



Model: PatchTST hybrid (72h)



Model: SimpleMOMENT

```
==== TEST CLASSIFICATION REPORT (will_fail_24h) ====
      precision    recall   f1-score   support
No Failure       0.998     0.935     0.965    126676
Failure          0.208     0.892     0.338     2424

accuracy          -         -         -     129100
macro avg        0.603     0.913     0.651    129100
weighted avg     0.983     0.934     0.954    129100

==== TEST CLASSIFICATION REPORT (will_fail_48h) ====
      precision    recall   f1-score   support
No Failure       0.981     0.956     0.968    124330
Failure          0.313     0.526     0.392     4770

accuracy          -         -         -     129100
macro avg        0.647     0.741     0.680    129100
weighted avg     0.957     0.940     0.947    129100

==== TEST CLASSIFICATION REPORT (will_fail_72h) ====
      precision    recall   f1-score   support
No Failure       0.957     0.972     0.965    122014
Failure          0.347     0.255     0.294     7086

accuracy          -         -         -     129100
macro avg        0.652     0.614     0.629    129100
weighted avg     0.924     0.933     0.928    129100
```

Implications and Future Work

- Near-perfect 24h prediction with PatchTST + contextual features (94% recall, 83% precision).
- Class imbalance challenge: failures = only 0.001% of data → rare event detection remains difficult.
- Long-term prediction hard: performance drops significantly beyond 24–48 hours.
- Traditional feature engineering still matters: raw sensor data alone not enough.
- Business value: improved recall/precision can save millions in downtime & false alarms.

Limitations:

- Couldn't run full MOMENT model due to hardware constraints.
- Transformers mainly built for forecasting, not classification → needed adaptation.

Future work:

- Handle imbalance with SMOTE, focal loss, or other methods.
- Try new transformer architectures (Informer, Autoformer, etc.).
- Improve computational efficiency for foundation models (e.g., full MOMENT).

Conclusion

- Transformers (esp. PatchTST + context) outperform traditional methods in short-term failure prediction.
- High recall + precision at 24h window = strong business impact.
- Challenges remain: long-term prediction + extreme imbalance.
- Transformers are evolving fast: currently forecasting-focused, but big potential for classification tasks.
- Predictive maintenance future: transformer architectures are a promising and scalable solution.

Q&A

References

- <https://iot-analytics.com/predictive-maintenance-market/>
- <https://www.kaggle.com/datasets/arnabbiswas1/microsoft-azure-predictive-maintenance>
- https://databricks.com/library/model/autonlab_moment-1-large/
- https://huggingface.co/docs/transformers/en/model_doc/patchtst
- <https://arxiv.org/html/2505.06295v1>
- https://www.researchgate.net/publication/345965691_PERFORMANCE_COMPARISON_OF_MACHINE_LEARNING_ALGORITHMS_FOR_PREDICTIVE_MAINTENANCE

Thank You!